

YARN

Apache Hadoop Next Generation Compute Platform



Apache Hadoop & YARN

- **Apache Hadoop**

- De facto Big Data open source platform
- Running for about 5 years in production at hundreds of companies like Yahoo, Ebay and Facebook

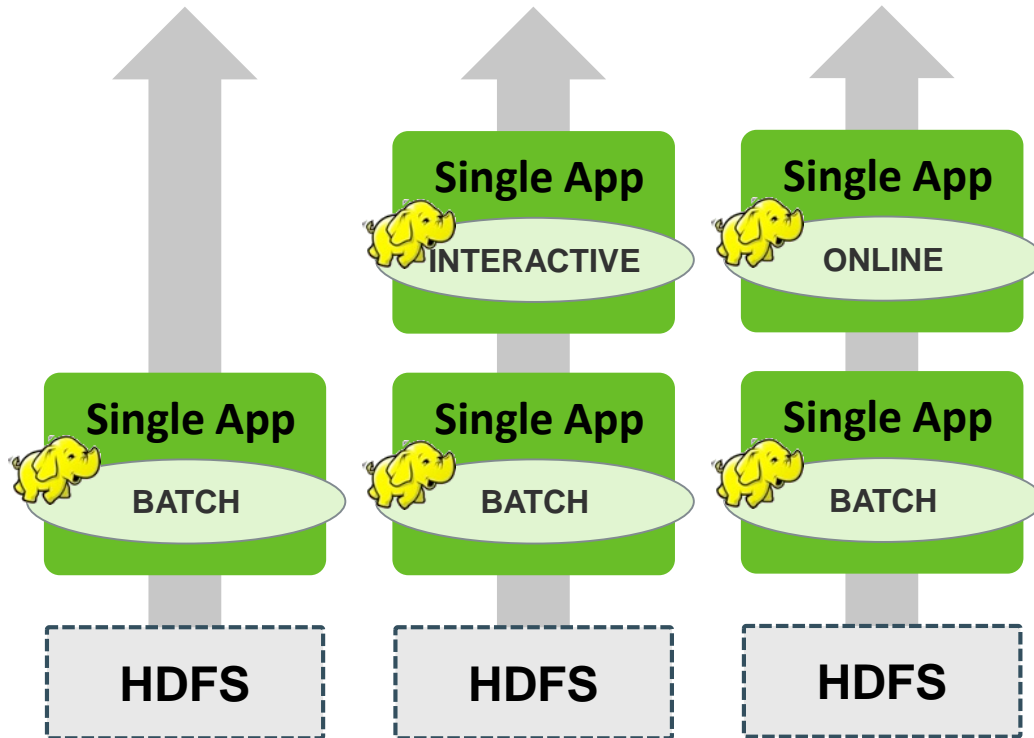
- **Hadoop 2**

- Significant improvements in HDFS distributed storage layer. High Availability, NFS, Snapshots
- YARN – next generation compute framework for Hadoop designed from the ground up based on experience gained from Hadoop 1
- YARN running in production at Yahoo for about a year
- YARN awarded Best Paper at SOCC 2013

1st Generation Hadoop: Batch Focus

HADOOP 1.0

Built for Web-Scale Batch Apps



All other usage patterns
MUST leverage same
infrastructure

**Forces Creation of Silos to
Manage Mixed Workloads**

Hadoop 1 Architecture


JobTracker

Manage Cluster Resources & Job Scheduling

TaskTracker

Per-node agent

Manage Tasks



Hadoop 1 Limitations

Lacks Support for Alternate Paradigms and Services

Force everything needs to look like Map Reduce

Iterative applications in MapReduce are 10x slower

Scalability

Max Cluster size ~5,000 nodes

Max concurrent tasks ~40,000

Availability

Failure Kills Queued & Running Jobs

Hard partition of resources into map and reduce slots

Non-optimal Resource Utilization

Our Vision: Hadoop as Next-Gen Platform

Single Use System

Batch Apps

HADOOP 1.0

MapReduce

(cluster resource management
& data processing)

HDFS

(redundant, reliable storage)



Multi Purpose Platform

Batch, Interactive, Online, Streaming, ...

HADOOP 2.0

MapReduce

(data processing)

Others

YARN

(cluster resource management)

HDFS2

(redundant, highly-available & reliable storage)

Hadoop 2 - YARN Architecture

ResourceManager (RM)

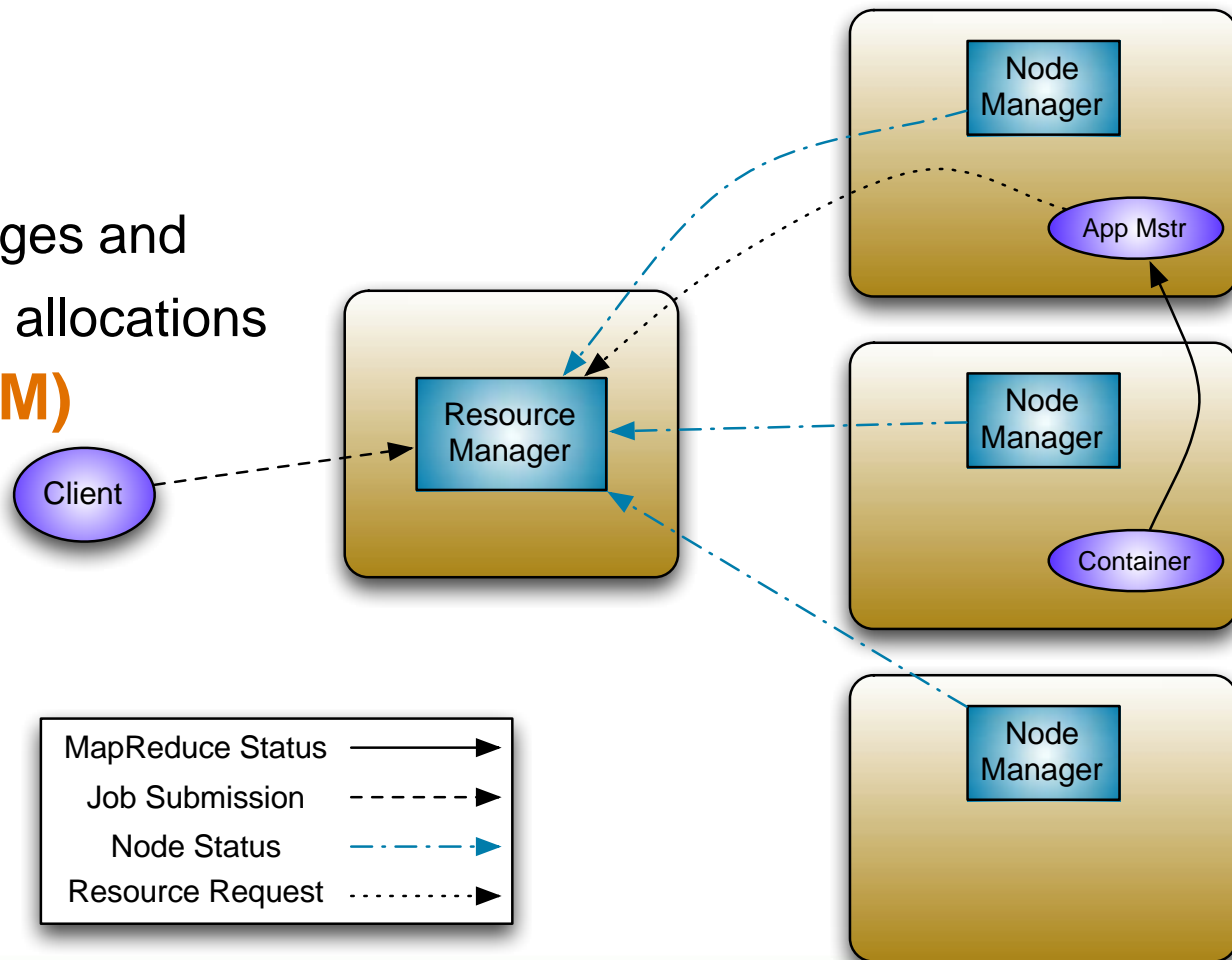
Central agent - Manages and allocates cluster resources

NodeManager (NM)

Per-Node agent - Manages and enforces node resource allocations

ApplicationMaster (AM)

Per-Application –
Manages application
lifecycle and task
scheduling

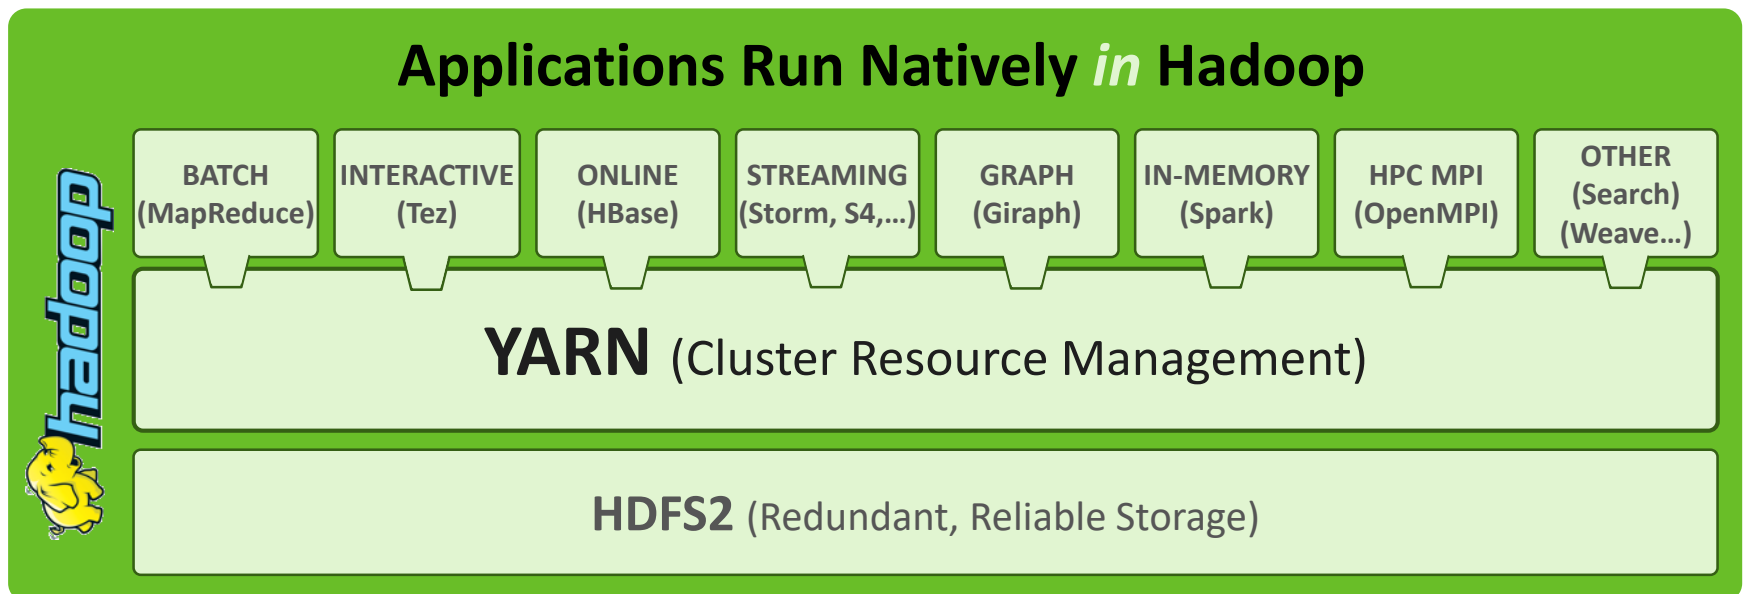


YARN: Taking Hadoop Beyond Batch

Store ALL DATA in one place...

Interact with that data in MULTIPLE WAYS

with Predictable Performance and Quality of Service



5 Key Benefits of YARN

- 1. New Applications & Services**
- 2. Improved cluster utilization**
- 3. Scale**
- 4. Experimental Agility**
- 5. Shared Services**



Key Improvements in YARN

Framework supporting multiple applications

- Separate generic resource brokering from application logic
- Define protocols/libraries and provide a framework for custom application development
- Share same Hadoop Cluster across applications

Cluster Utilization

- Generic resource container model replaces fixed Map/Reduce slots. Container allocations based on locality, memory (CPU coming soon)
- Sharing cluster among multiple application

Key Improvements in YARN

Scalability

- Removed complex app logic from RM, scale further
- State machine, message passing based loosely coupled design
- Compact scheduling protocol

Application Agility and Innovation

- Use Protocol Buffers for RPC gives wire compatibility
- Map Reduce becomes an application in user space unlocking safe innovation
- Multiple versions of an app can co-exist leading to experimentation
- Easier upgrade of framework and application

Key Improvements in YARN

Shared Services

- Common services needed to build distributed application are included in a pluggable framework
- Distributed file sharing service
- Remote data read service
- Log Aggregation Service

YARN: Efficiency with Shared Services



Yahoo! leverages YARN

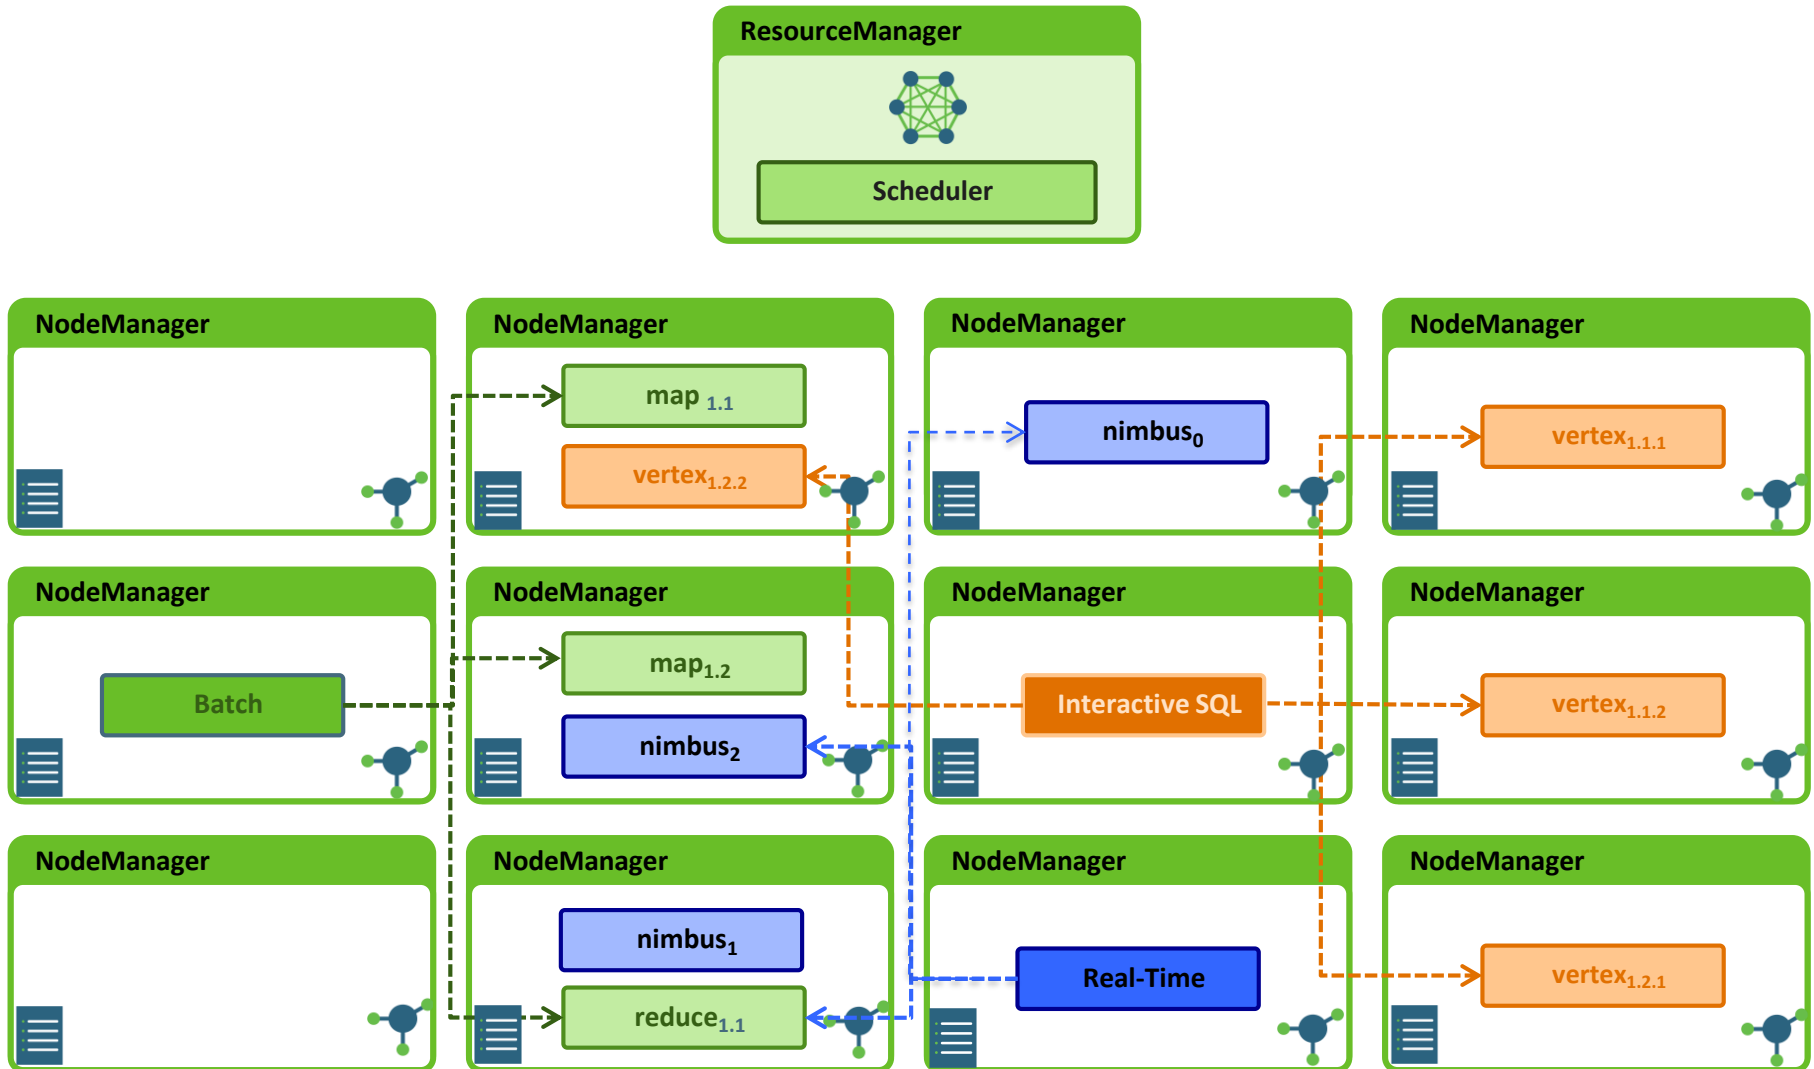
40,000+ nodes running YARN across over 365PB of data
~400,000 jobs per day for about 10 million hours of compute time

Estimated a 60% – 150% improvement on node usage per day using YARN

Eliminated Colo (~10K nodes) due to increased utilization

For more details check out the YARN SOCC 2013 paper

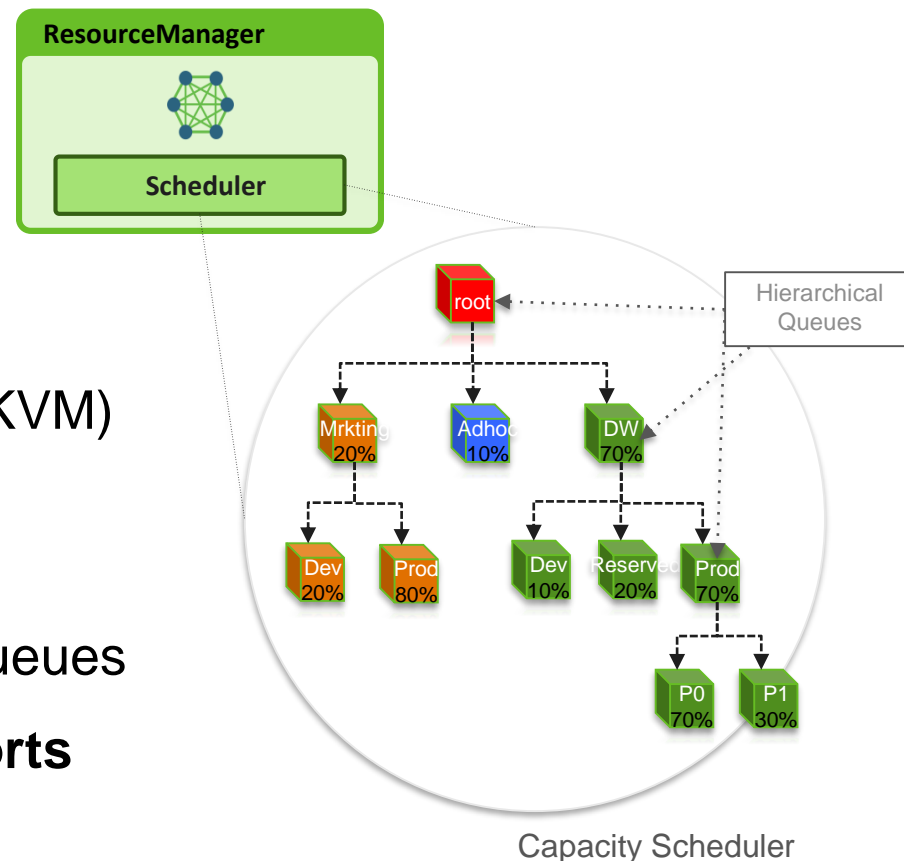
YARN as Cluster Operating System



Multi-Tenancy is Built-in

- **Queues**
- **Economics as *queue-capacity***
 - Hierarchical Queues
- **SLAs**
 - Cooperative Preemption
- **Resource Isolation**
 - Linux: cgroups
 - Roadmap: Virtualization (Xen, KVM)
- **Administration**
 - Queue ACLs
 - Run-time re-configuration for queues

Default Capacity Scheduler supports all features



YARN Eco-system

Applications Powered by YARN

Apache Giraph – Graph Processing

Apache Hama - BSP

Apache Hadoop MapReduce – Batch

Apache Tez – Batch/Interactive

Apache S4 – Stream Processing

Apache Samza – Stream Processing

Apache Storm – Stream Processing

Apache Spark – Iterative applications

Elastic Search – Scalable Search

Cloudera Llama – Impala on YARN

DataTorrent – Data Analysis

HOYA – HBase on YARN



**There's an app for that...
YARN App Marketplace!**

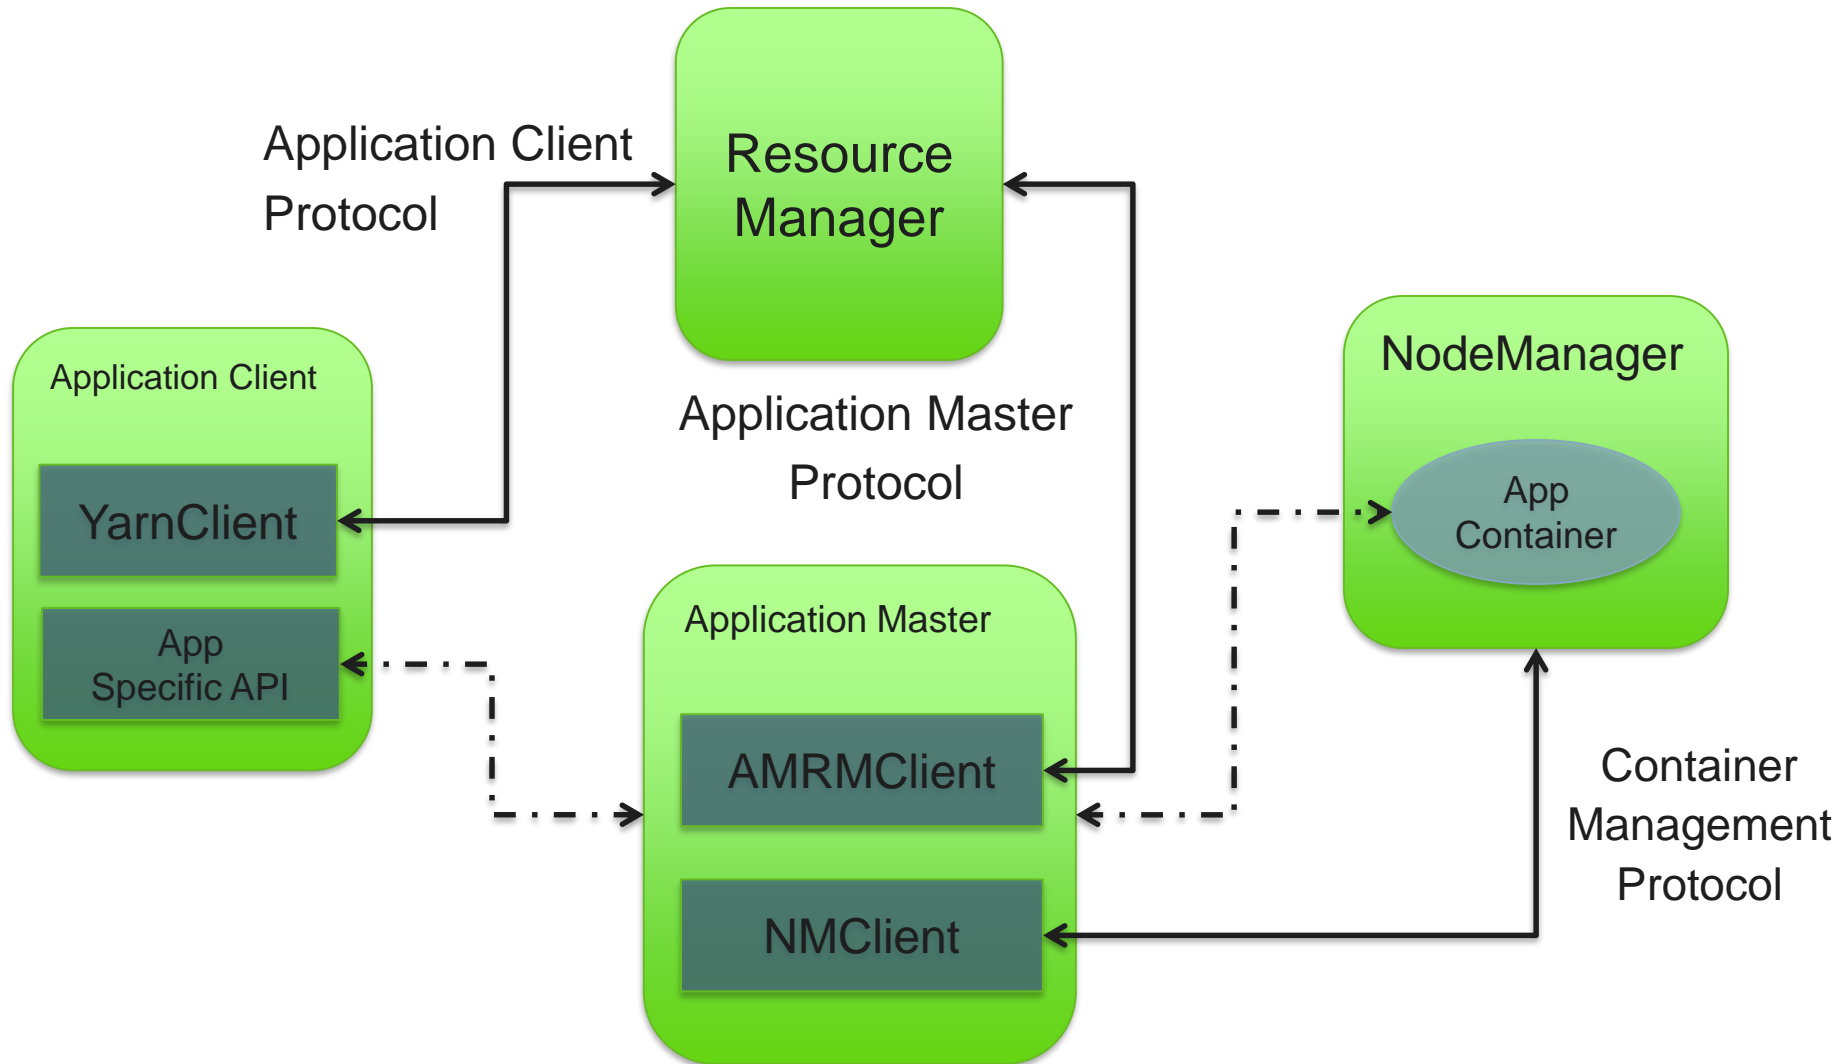
Frameworks Powered By YARN

Apache Twill

REEF by Microsoft

Spring support for Hadoop 2

YARN Application Lifecycle



BYOA – Bring Your Own App

Application Client Protocol: Client to RM interaction

- Library: YarnClient
- Application Lifecycle control
- Access Cluster Information

Application Master Protocol: AM – RM interaction

- Library: AMRMClient / AMRMClientAsync
- Resource negotiation
- Heartbeat to the RM

Container Management Protocol: AM to NM interaction

- Library: NMClient/NMClientAsync
- Launching allocated containers
- Stop Running containers

Use external frameworks like Twill/REEF/Spring

YARN Future Work

- **ResourceManager High Availability**
 - Automatic failover
 - Work preserving failover
- **Scheduler Enhancements**
 - SLA Driven Scheduling, Low latency allocations
 - Multiple resource types – disk/network/GPUs/affinity
- **Rolling upgrades**
- **Generic History Service**
- **Long running services**
 - Better support to running services like HBase
 - Service Discovery
- **More utilities/libraries for Application Developers**
 - Failover/Checkpointing

Key Take-Aways

- **YARN is a platform to build/run Multiple Distributed Applications in Hadoop**
- **YARN is completely Backwards Compatible for existing MapReduce apps**
- **YARN enables Fine Grained Resource Management via Generic Resource Containers.**
- **YARN has built-in support for multi-tenancy to share cluster resources and increase cost efficiency**
- **YARN provides a cluster operating system like abstraction for a modern data architecture**

Apache YARN

The Data Operating System for Hadoop 2.0

Flexible

Enables other purpose-built data processing models beyond MapReduce (batch), such as interactive and streaming

Efficient

Increase processing **IN** Hadoop on the same hardware while providing predictable performance & quality of service

Shared

Provides a stable, reliable, secure foundation and shared operational services across multiple workloads

Data Processing Engines Run Natively **IN** Hadoop

BATCH
MapReduce

INTERACTIVE
Tez

ONLINE
HBase

STREAMING
Storm, S4, ...

GRAPH
Giraph

MICROSOFT
REEF

SAS
LASR, HPA

OTHERS

YARN: Cluster Resource Management

HDFS2: Redundant, Reliable Storage



Thank you!



Download Sandbox: Experience Apache Hadoop
Both 2.0 and 1.x Versions Available!

<http://hortonworks.com/products/hortonworks-sandbox/>

Questions?