

Big Data and Cloud Computing: Current State and Future Opportunities

Why? WEB is replacing the Desktop



Paradigm Shift in Computing

Azure Services Platform



What is Cloud Computing?

- Delivering applications and services over the Internet:
 - Software as a service
- Extended to:
 - Infrastructure as a service: Amazon EC2
 - Platform as a service: Google AppEngine, Microsoft Azure
- Utility Computing: pay-as-you-go computing
 - Illusion of infinite resources
 - No up-front cost
 - Fine-grained billing (e.g. hourly)

Cloud Computing: History

“ If computers of the kind I have advocated become the computers of the future, then computing may someday be organized as a public utility just as the telephone system is a public utility... The computer utility could become the basis of a new and important industry. ”

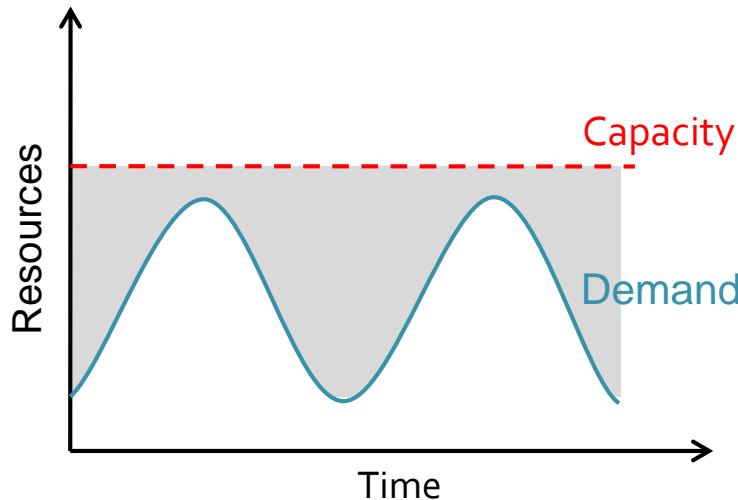
—John McCarthy, speaking at the MIT Centennial in 1961^[2]

Cloud Computing: Why Now?

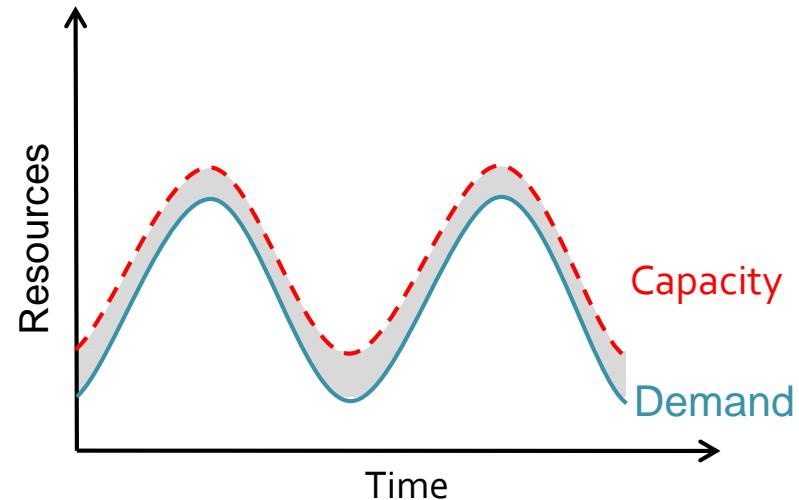
- Experience with very large datacenters
 - Unprecedented economies of scale
 - Transfer of risk
- Technology factors
 - Pervasive broadband Internet
 - Maturity in Virtualization Technology
- Business factors
 - Minimal capital expenditure
 - Pay-as-you-go billing model

Economics of Cloud Users

- Pay by use instead of provisioning for peak



Static data center



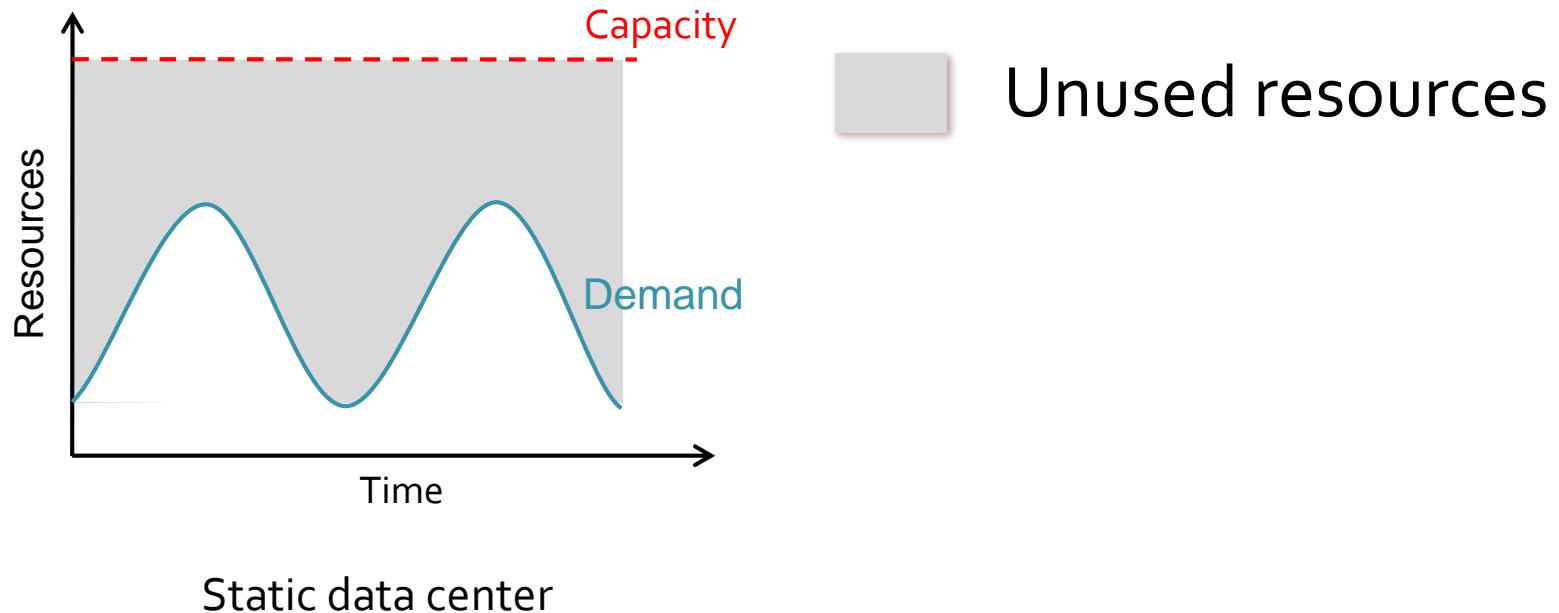
Data center in the cloud



Unused resources

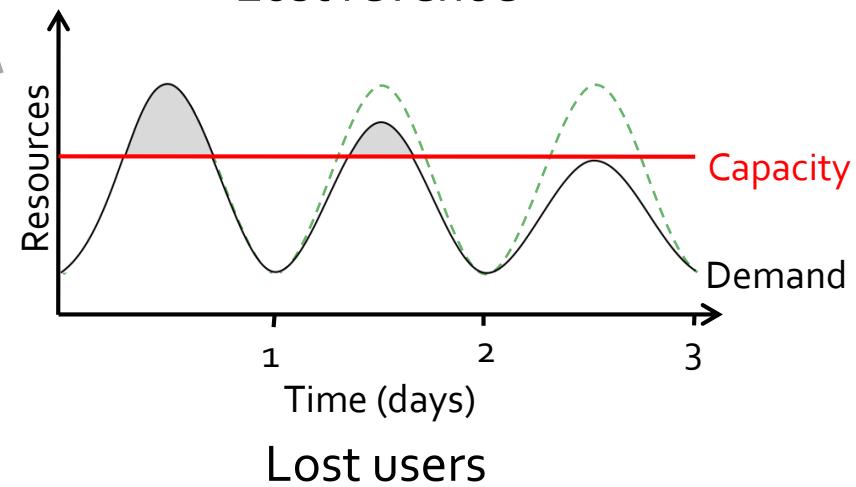
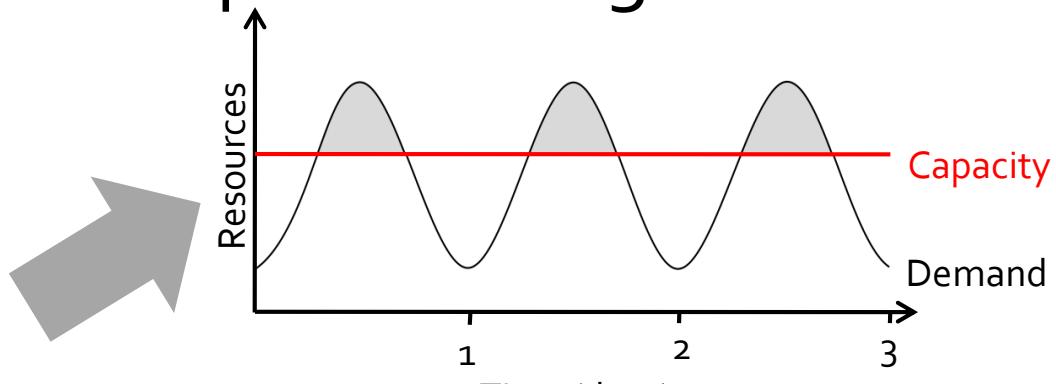
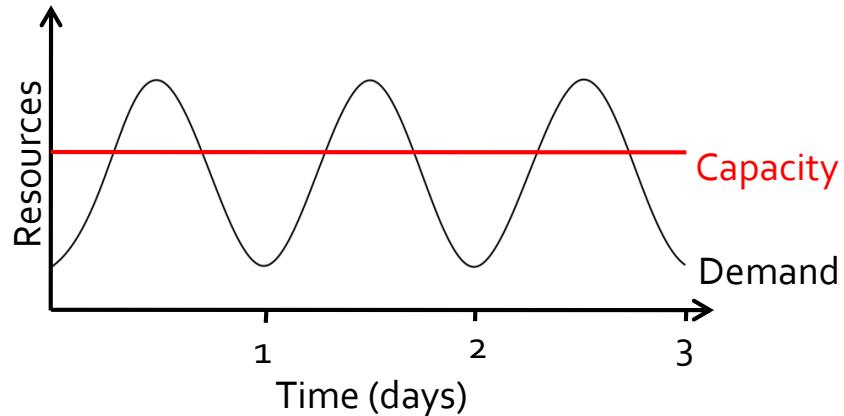
Economics of Cloud Users

- Risk of over-provisioning: underutilization



Economics of Cloud Users

- Heavy penalty for under-provisioning



The Big Picture

- Unlike the earlier attempts:
 - Distributed Computing
 - Distributed Databases
 - Grid Computing
- Cloud Computing is likely to persist:
 - Organic growth: Google, Yahoo, Microsoft, and Amazon
 - Poised to be an integral aspect of National Infrastructure in US and other countries

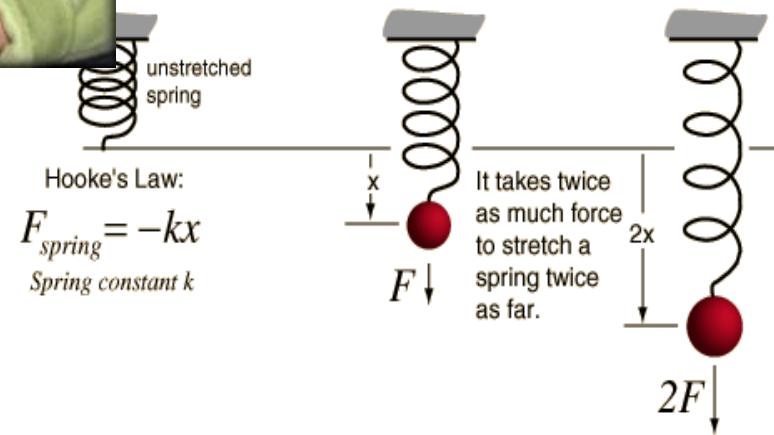
Cloud Reality

- Facebook Generation of Application Developers
- Animoto.com:
 - Started with 50 servers on Amazon EC2
 - Growth of 25,000 users/hour
 - Needed to scale to 3,500 servers in 2 days
(RightScale@SantaBarbara)
- Many similar stories:
 - RightScale
 - Joyent
 - ...

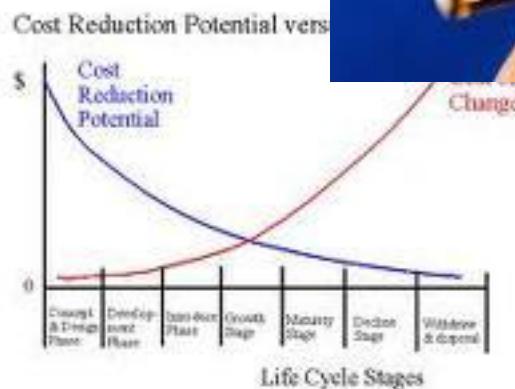
Cloud Challenges: Elasticity



Big Data 2019 Spring



Cloud Challenges: Differential Pricing Models



* Adapted from Cusumano Exhibit 2, p. 15



Outline

- **Data in the Cloud**
 - Platforms for Data Analysis
 - Platforms for Update intensive workloads
- Data Platforms for Large Applications
- Multitenant Data Platforms
- Open Research Challenges

Our Data-driven World

- Science
 - Data bases from astronomy, genomics, environmental data, transportation data, ...
- Humanities and Social Sciences
 - Scanned books, historical documents, social interactions data, ...
- Business & Commerce
 - Corporate sales, stock market transactions, census, airline traffic, ...
- Entertainment
 - Internet images, Hollywood movies, MP3 files, ...
- Medicine
 - MRI & CT scans, patient records, ...

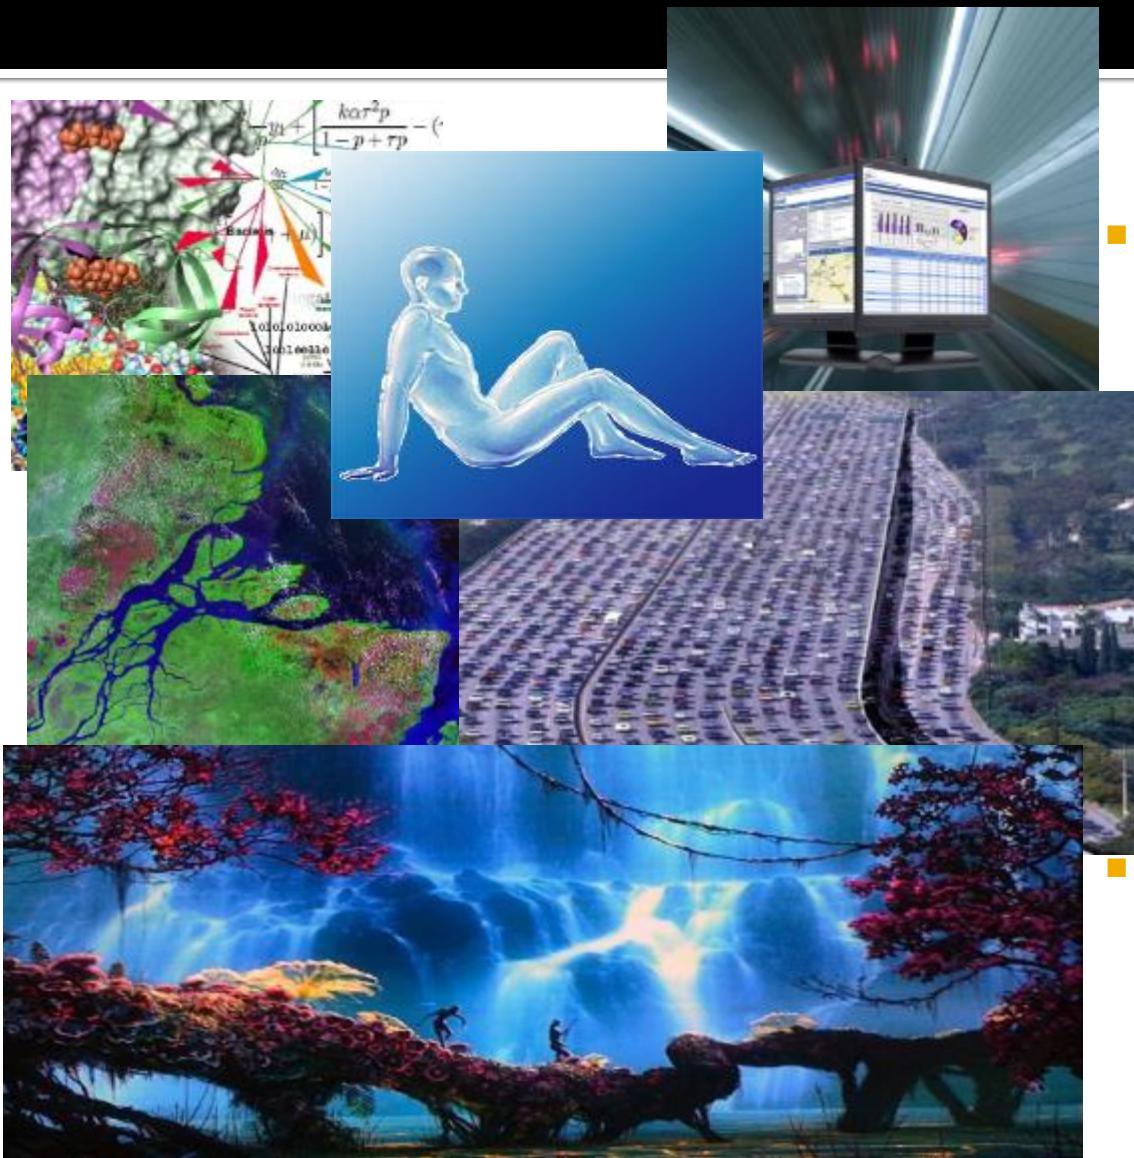
Data-rich World

- Data capture and collection:
 - Highly instrumented environment
 - Sensors and Smart Devices
 - Network

- Data storage:
 - Seagate 1 TB Barracuda @ \$72.95 from Amazon.com (73¢/GB)



What can we do with this wealth?



- What can we do?
 - Scientific breakthroughs
 - Business process efficiencies
 - Realistic special effects
 - Improve quality-of-life: healthcare, transportation, environmental disasters, daily life, ...
- Could We Do More?
 - YES: but need major advances in our capability to analyze this data

Cloud Computing Modalities



“Can we outsource our IT software and hardware infrastructure?”

- Hosted Applications and services
- Pay-as-you-go model
- Scalability, fault-tolerance, elasticity, and self-manageability



“We have terabytes of click-stream data – what can we do with it?”

- Very large data repositories
- Complex analysis
- Distributed and parallel data processing

Outline

- **Data in the Cloud**
 - **Platforms for Data Analysis**
 - Platforms for Update intensive workloads
- Data Platforms for Large Applications
- Multitenant Data Platforms
- Open Research Challenges

Data Warehousing, Data Analytics & Decision Support Systems

- Used to manage and control business
- Transactional Data: historical or point-in-time
- Optimized for inquiry rather than update
- Use of the system is loosely defined and can be ad-hoc
- Used by managers and analysts to understand the business and make judgments

Data Analytics in the Web Context

- Data capture at the user interaction level:
 - in contrast to the client transaction level in the Enterprise context
- As a consequence the amount of data increases significantly
- Greater need to analyze such data to understand user behaviors

Data Analytics in the Cloud

- Scalability to large data volumes:
 - Scan 100 TB on 1 node @ 50 MB/sec = 23 days
 - Scan on 1000-node cluster = 33 minutes
- Divide-And-Conquer (i.e., data partitioning)
- Cost-efficiency:
 - Commodity nodes (cheap, but unreliable)
 - Commodity network
 - Automatic fault-tolerance (fewer administrators)
 - Easy to use (fewer programmers)

Platforms for Large-scale Data Analysis

- **Parallel DBMS technologies**
 - Proposed in the late eighties
 - Matured over the last two decades
 - Multi-billion dollar industry: Proprietary DBMS Engines intended as Data Warehousing solutions for very large enterprises
- **Map Reduce**
 - pioneered by Google
 - popularized by Yahoo! (Hadoop)

Parallel DBMS technologies

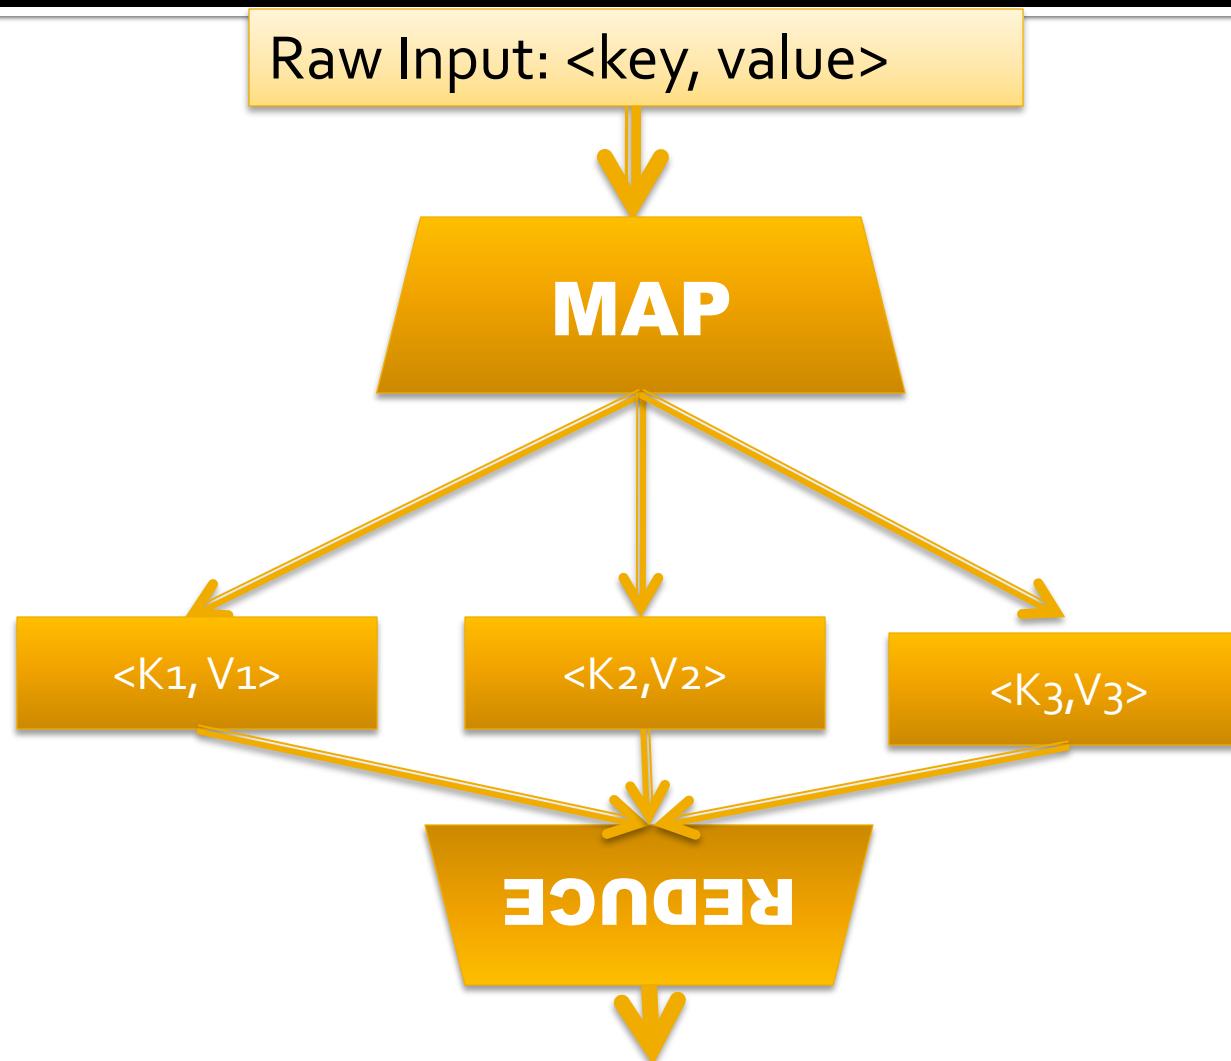
- Popularly used for more than two decades
 - Research Projects: Gamma, Grace, ...
 - Commercial: Multi-billion dollar industry but access to only a privileged few
- Relational Data Model
- Indexing
- Familiar SQL interface
- Advanced query optimization
- *Well understood and well studied*

MapReduce

[Dean et al., OSDI 2004, CACM Jan 2008, CACM Jan 2010]

- Overview:
 - Data-parallel programming model
 - An associated parallel and distributed implementation for commodity clusters
- Pioneered by Google
 - Processes 20 PB of data per day
- Popularized by open-source Hadoop project
 - Used by Yahoo!, Facebook, Amazon, and the list is growing ...

Programming Framework



MapReduce Advantages

- Automatic Parallelization:
 - Depending on the size of RAW INPUT DATA → instantiate multiple MAP tasks
 - Similarly, depending upon the number of intermediate <key, value> partitions → instantiate multiple REDUCE tasks
- Run-time:
 - Data partitioning
 - Task scheduling
 - Handling machine failures
 - Managing inter-machine communication
- Completely transparent to the programmer/analyst/user

MapReduce Experience

- Runs on large commodity clusters:
 - 1000s to 10,000s of machines
- Processes many terabytes of data
- Easy to use since run-time complexity hidden from the users
- 1000s of MR jobs/day at Google (circa 2004)
- 100s of MR programs implemented (circa 2004)

The Need

- Special-purpose programs to process large amounts of data: crawled documents, Web Query Logs, etc.
- At Google and others (Yahoo!, Facebook):
 - Inverted index
 - Graph structure of the WEB documents
 - Summaries of #pages/host, set of frequent queries, etc.
 - Ad Optimization
 - Spam filtering

MapReduce Contributions

**Simple & Powerful
Programming Paradigm
For
Large-scale Data Analysis**

**Run-time System
For
Large-scale Parallelism &
Distribution**

Takeaway

- MapReduce's data-parallel programming model hides complexity of distribution and fault tolerance
- Key philosophy:
 - *Make it scale*, so you can throw hardware at problems
 - *Make it cheap*, saving hardware, programmer and administration costs (but requiring fault tolerance)
- Hive and Pig further simplify programming
- MapReduce is not suitable for all problems, but when it works, it may save you a lot of time

Map Reduce vs Parallel DBMS

[Pavlo et al., SIGMOD 2009, Stonebraker et al., CACM 2010, ...]

	Parallel DBMS	MapReduce
Schema Support	✓	Not out of the box
Indexing	✓	Not out of the box
Programming Model	Declarative (SQL)	Imperative (C/C++, Java, ...) Extensions through Pig and Hive
Optimizations (Compression, Query Optimization)	✓	Not out of the box
Flexibility	Not out of the box	✓
Fault Tolerance	Coarse grained techniques	✓

MapReduce: A step backwards?

- Don't need 1000 nodes to process petabytes:
 - Parallel DBs do it in fewer than 100 nodes
- No support for schema:
 - Sharing across multiple MR programs difficult
- No indexing:
 - Wasteful access to unnecessary data
- Non-declarative programming model:
 - Requires highly-skilled programmers
- No support for JOINs:
 - Requires multiple MR phases for the analysis

MapReduce VS Parallel DB: Our 2¢

- Web application data is inherently distributed on a large number of sites:
 - Funneling data to DB nodes is a failed strategy
- Distributed and parallel programs difficult to develop:
 - Failures and dynamics in the cloud
- Indexing:
 - Sequential Disk access 10 times faster than random access.
 - Not clear if indexing is the right strategy.
- Complex queries:
 - DB community needs to JOIN hands with MR

Leveraging DBMS Technology for Scalable Analytics

Scalable Analytics Innovations

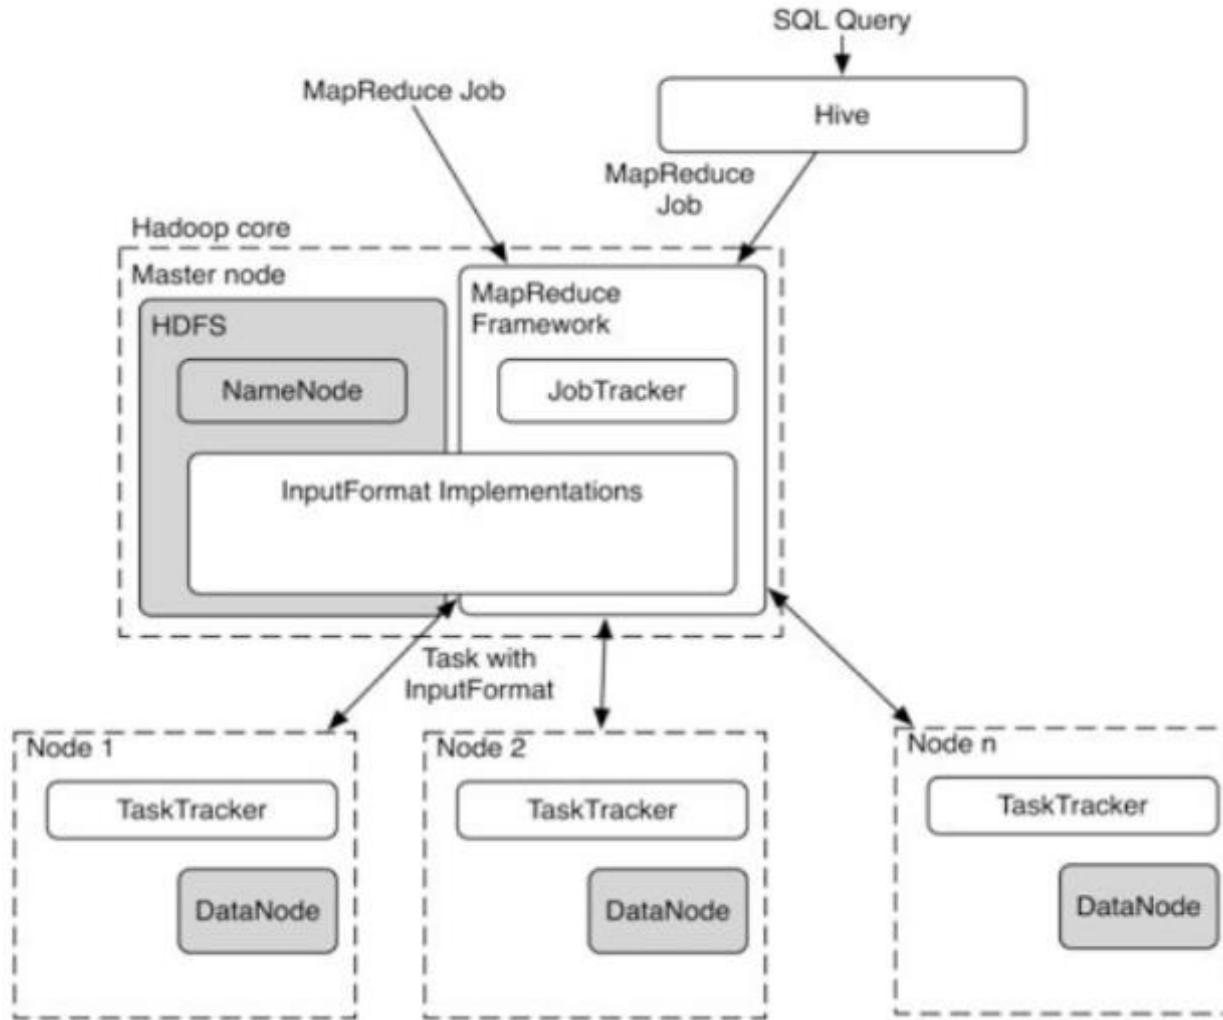
- Asynchronous Views over Cloud Data
 - Yahoo! Research (SIGMOD'2009)
- DataPath: Data-centric Analytic Engine
 - U Florida (Dobra) & Rice (Jermaine) (SIGMOD'2010)
- MapReduce innovations:
 - MapReduce Online (UCBerkeley)
 - HadoopDB (Yale)
 - Multi-way Join in MapReduce (Afrati&Ullman: EDBT'2010)
 - Hadoop++ (Dittrich et al.: VLDB'2010) and others

Hadoop DB – A Hybrid Approach

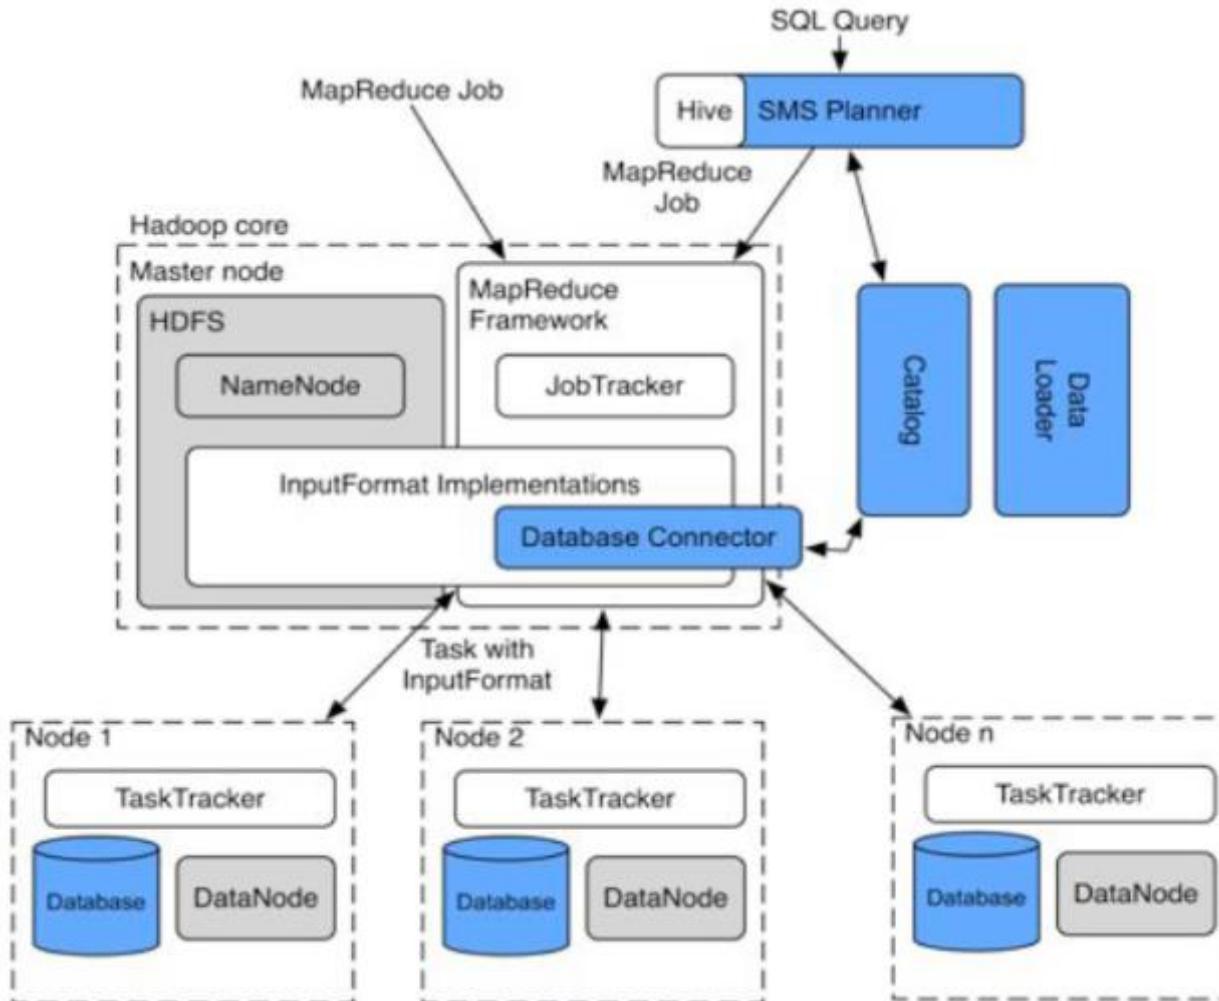
[Abouzeid et al., VLDB 2009]

- An architectural hybrid of MapReduce and DBMS technologies
- Use Fault-tolerance and Scale of MapReduce framework like Hadoop
- Leverage advanced data processing techniques of an RDBMS
- Expose a declarative interface to the user
- ***Goal: Leverage from the best of both worlds***

Architecture of HadoopDB



Architecture of HadoopDB



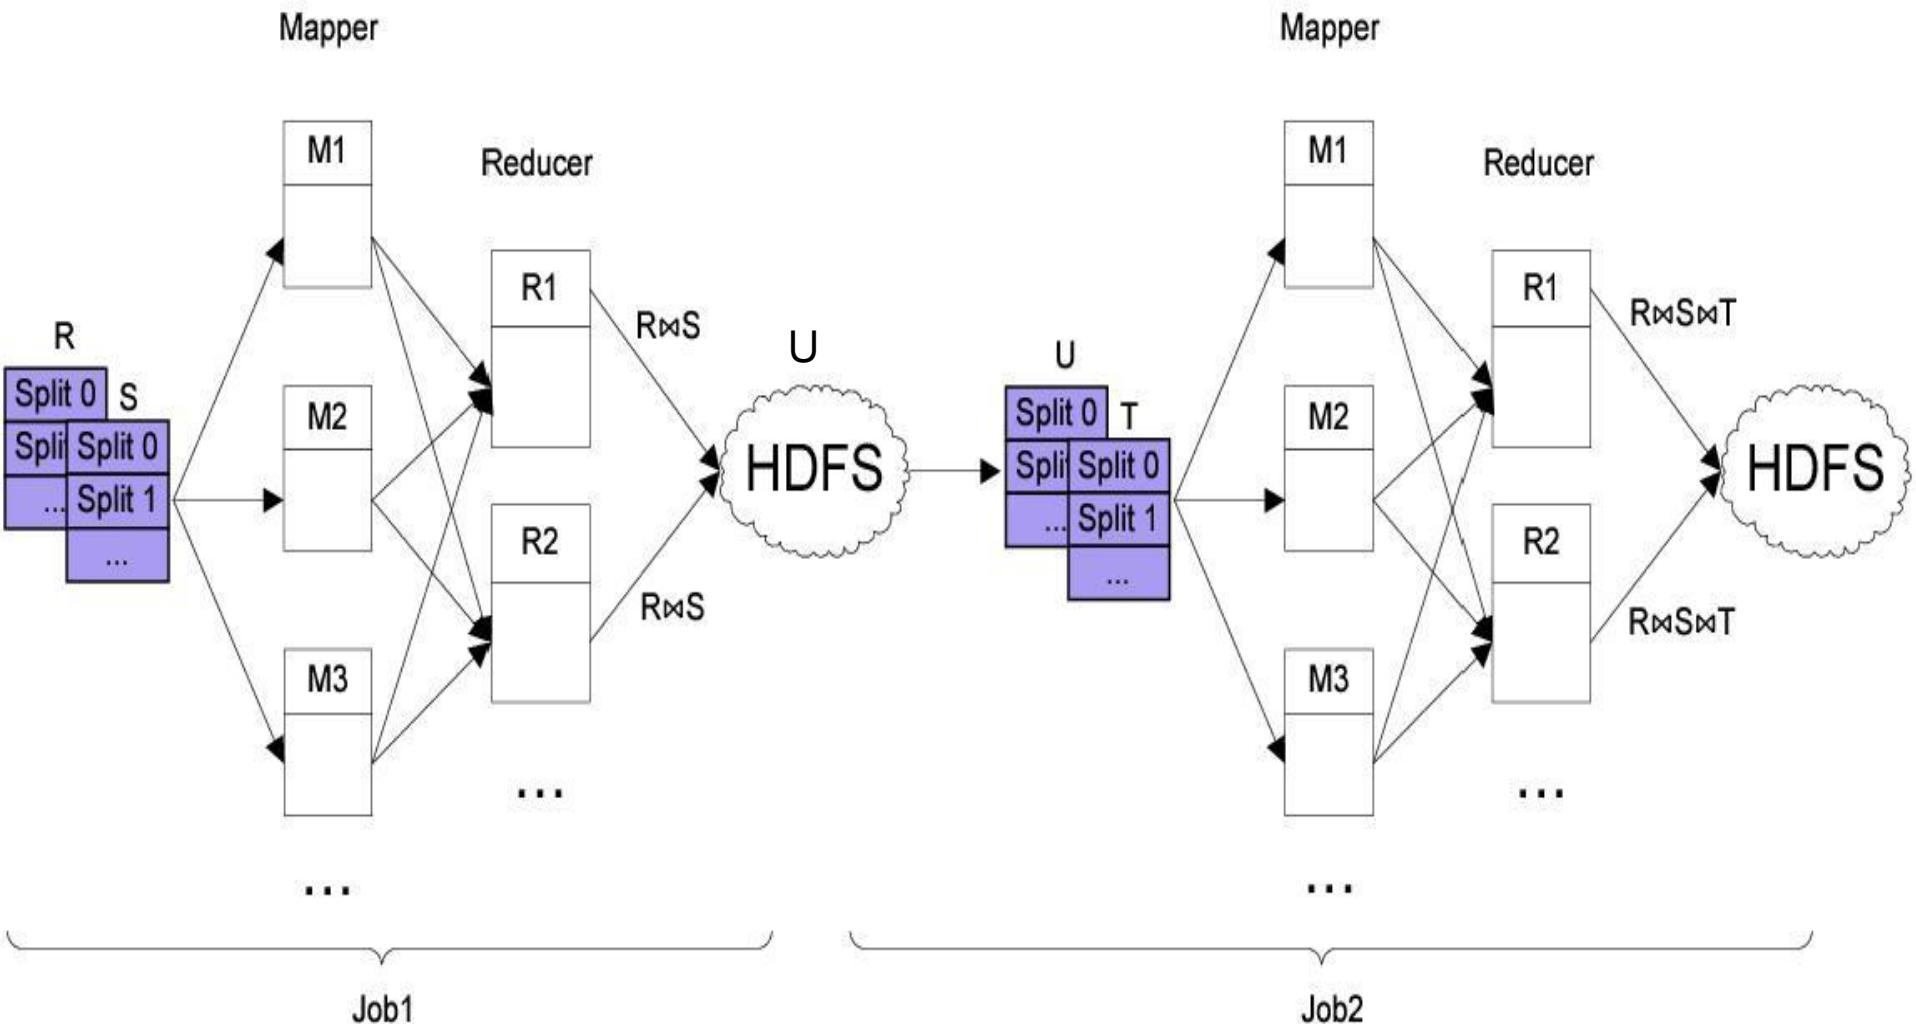
Two-way Joins and MapReduce

- MapReduce works with a single data source (table):
 - $\langle \text{key } K, \text{ value } V \rangle$
- How to use the MR framework to compute:
 - $R(A, B) \bowtie S(B, C)$
- Simple extension (proposed independently by multiple researchers):
 - $\langle a, b \rangle$ from R is mapped as: $\langle b, [R, a] \rangle$
 - $\langle b, c \rangle$ from S is mapped as: $\langle b, [S, c] \rangle$
- During the reduce phase:
 - Join the key-value pairs with the same key but different relations

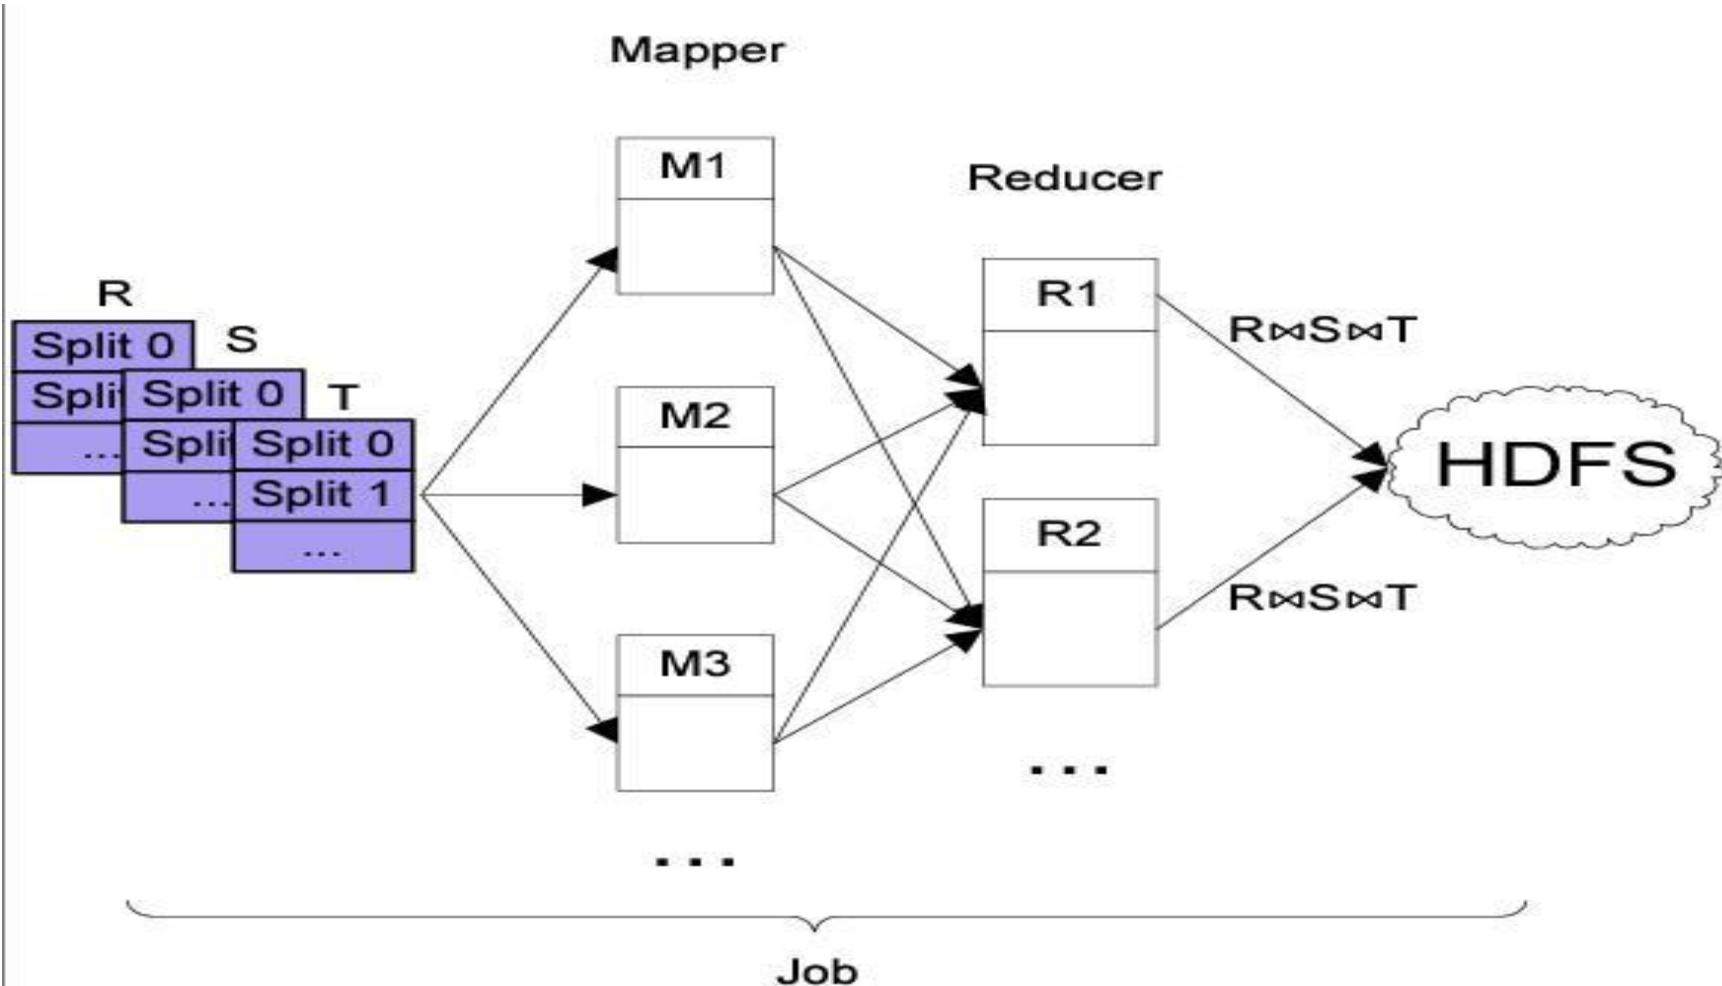
Beyond Two-way Joins?

- How to generalize to:
 - $R(A, B) \bowtie S(B, C) \bowtie T(C, D)$ in MapReduce?

Beyond Two-way Joins

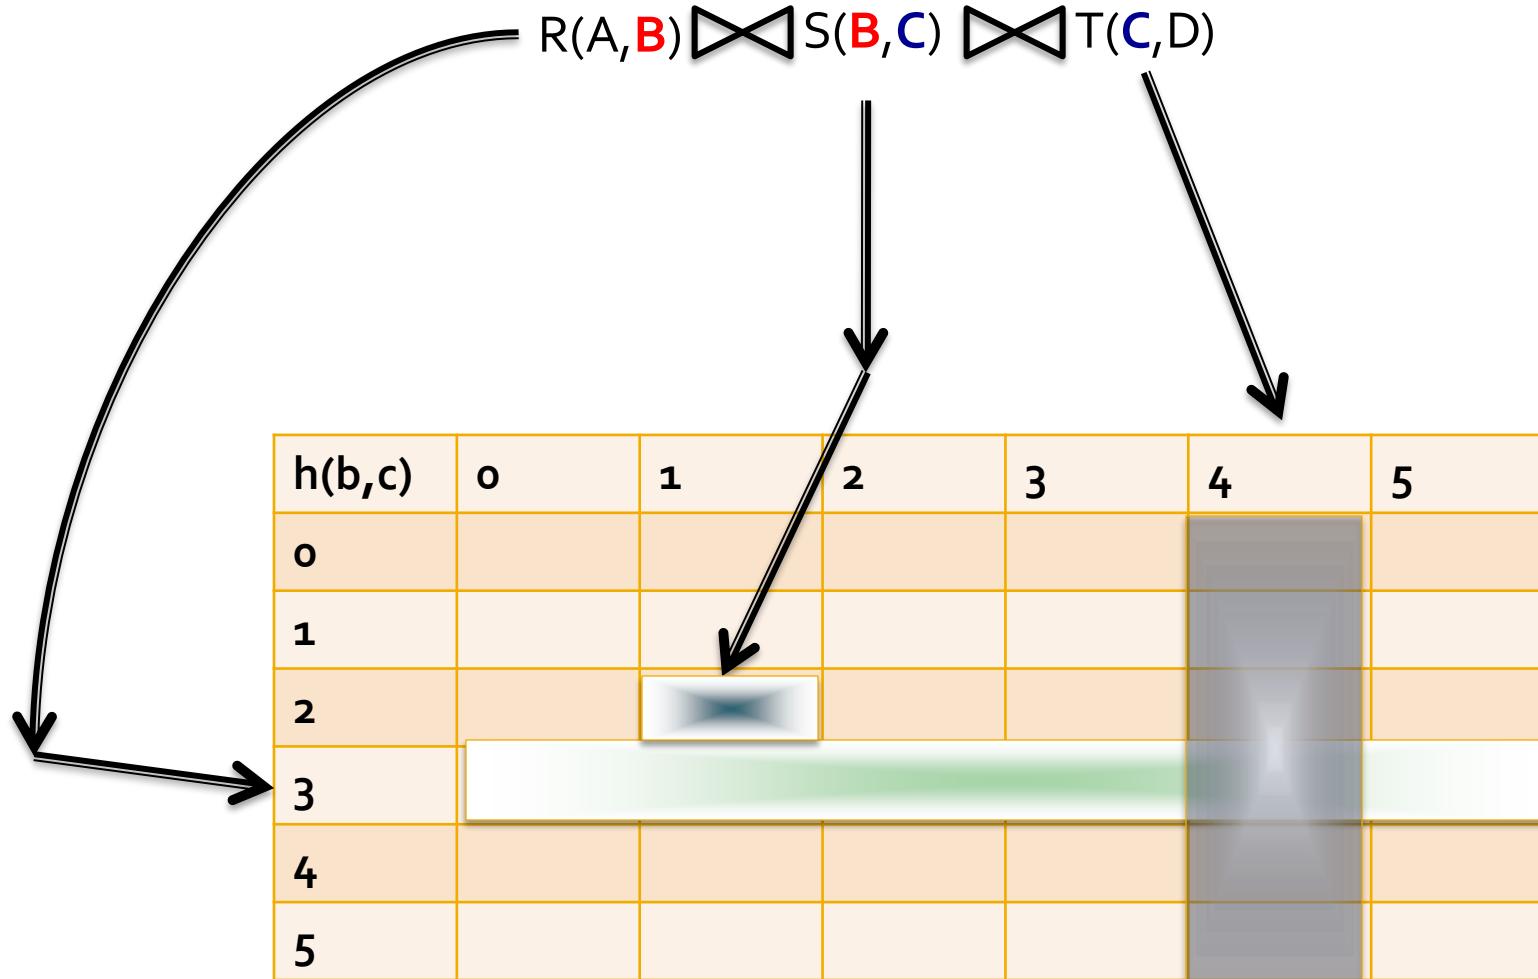


Beyond Two-way Joins?



Multi-way Join

[Afrati & Ullman: EDBT'2010]



OLAP over MapReduce

- Multi-way Join over a single MR phase:
 - One-to-many **shuffle**
- Communication cost:
 - 3-way Join: $O(n)$ communication
 - 4-way Join: $O(n^2)$
 - ...
 - M-way Join: $O(n^{M-2})$
- Clearly, not feasible for OLAP:
 - Large number of Dimension Tables
 - Many OLAP queries involve Join of FACT table with multiple Dimension tables.



Scalable OLAP

[SIGMOD'2011: NUS & UCSB Collaboration]

- Observations:
 - STAR Schema (or its variant)
 - DIMENSION tables typically are not very large (in most cases)
 - FACT table, on the other hand, have a large number of rows.
- Design approach:
 - Use MapReduce as a distributed & parallel computing substrate
 - Fact tables partitioned across multiple nodes
 - Partitioning strategy should be based on the dimension that are most often used as selection constraint in DW queries
 - Dimension tables are replicated
 - MAP tasks perform the STAR joins
 - REDUCE tasks then perform the aggregation

Data Analytics in the Cloud

New Challenges and Opportunities

Looking Forward

New Applications

- Complex data processing – Graphs and beyond
- Multidimensional Data Analytics: Location-based data
- Physical and Virtual Worlds: Social Networks and Social Media data & analysis

Conjectures

- New breed of Analysts:
 - Information-savvy users
 - Most users will become nimble analysts
 - Most transactional decisions will be preceded by a detailed analysis
- Convergence of OLAP and OLTP:
 - Both from the application point-of-view and from the infrastructure point-of-view

Outline

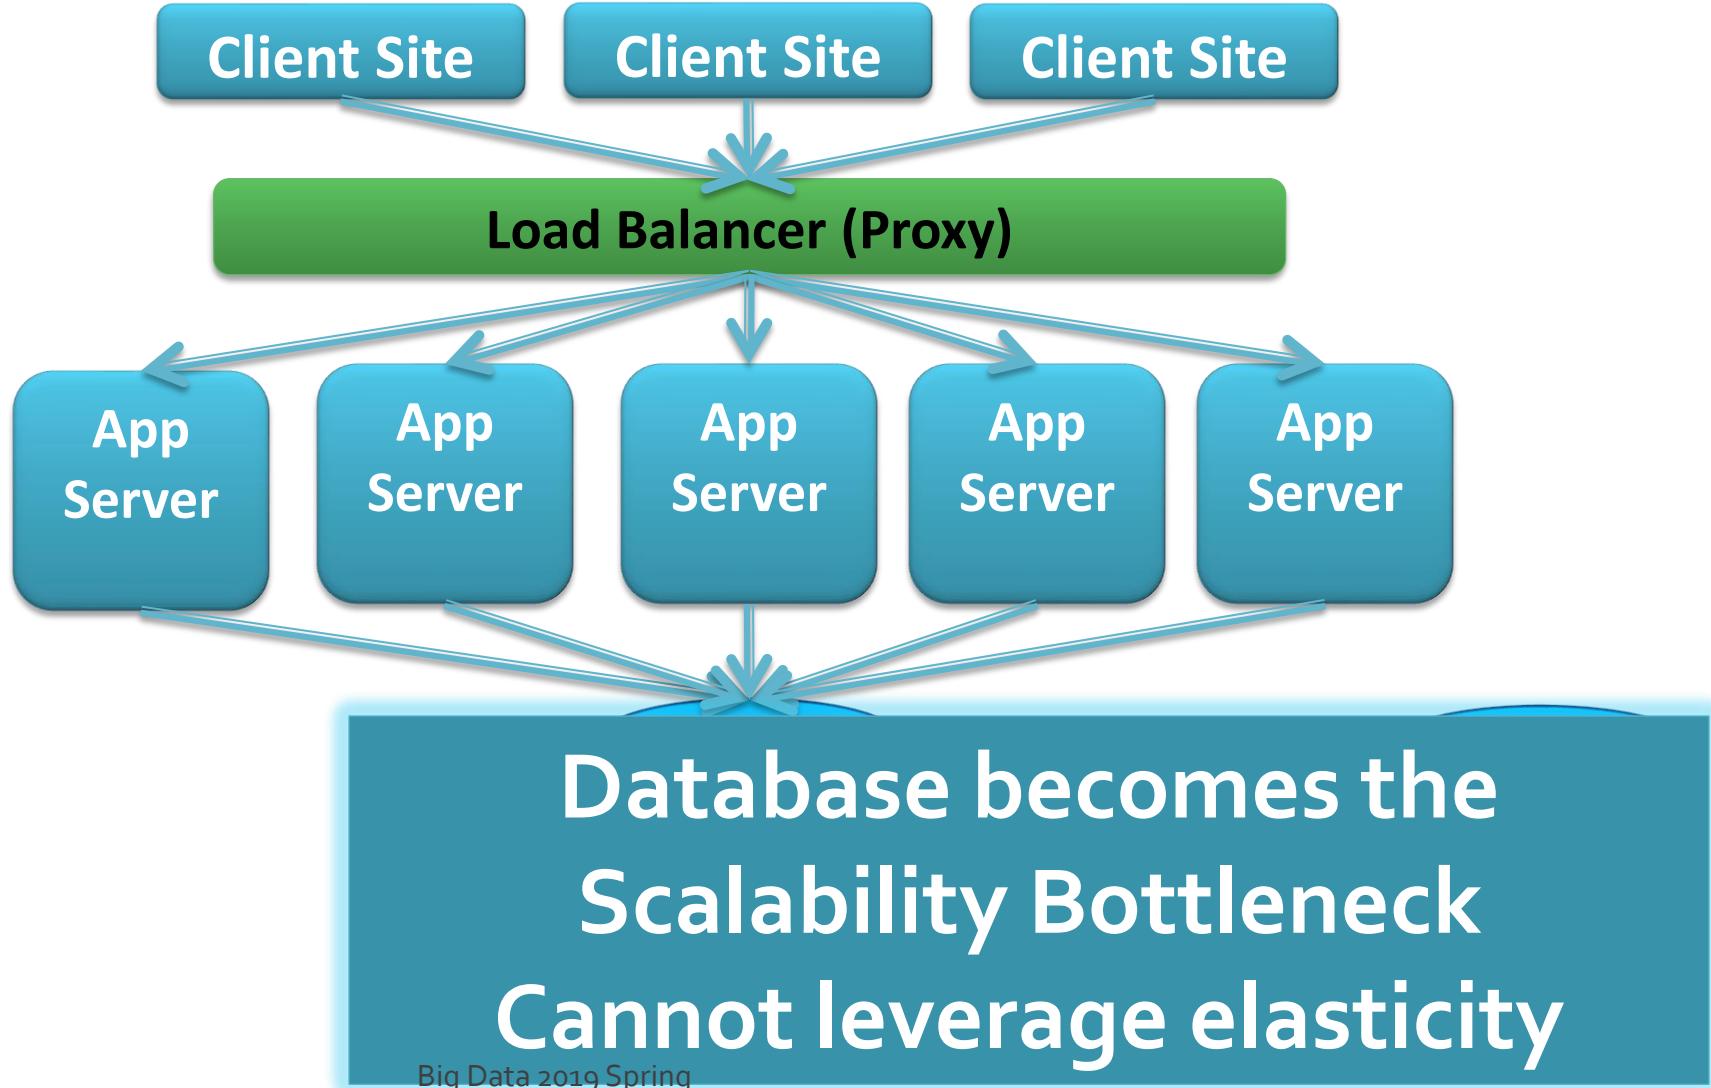
- **Data in the Cloud**
 - Platforms for Data Analysis
 - **Platforms for Update intensive workloads**
- Data Platforms for Large Applications
- Multitenant Data Platforms
- Open Research Challenges

Platforms for Update intensive workloads

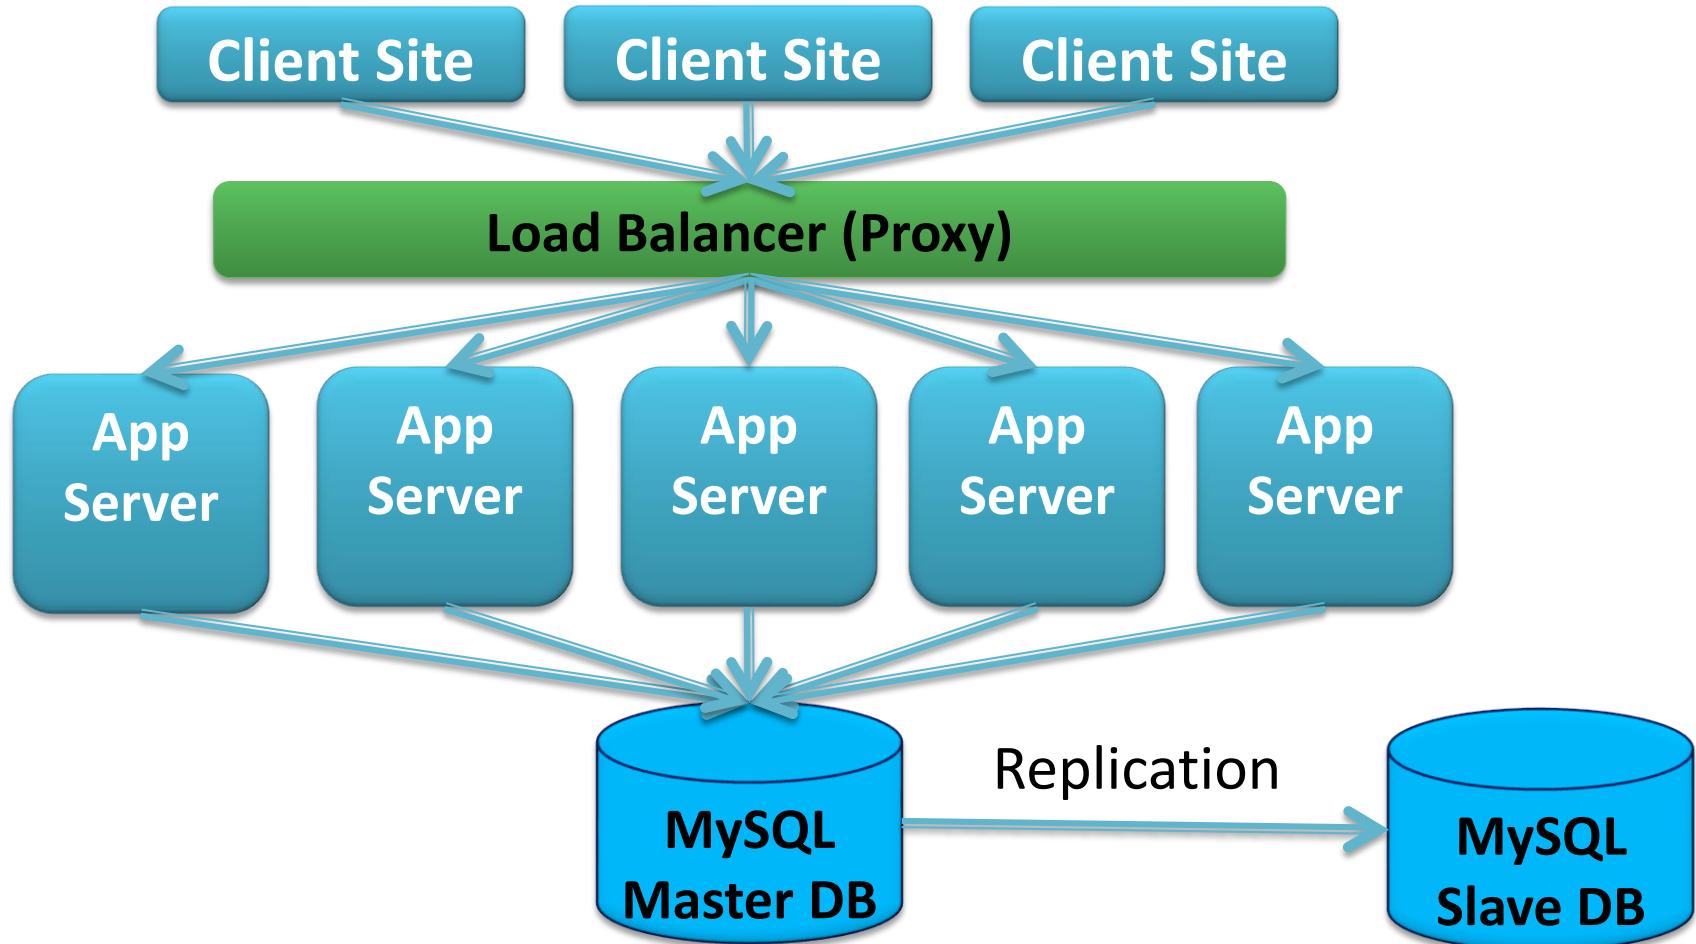
- Most enterprise solutions are based on RDBMS technology.
- Significant Operational Challenges:
 - Provisioning for Peak Demand
 - Resource under-utilization
 - Capacity planning: too many variables
 - Storage management: a massive challenge
 - System upgrades: extremely time-consuming
 - Complex mine-field of software and hardware licensing

→ Unproductive use of people-resources from a company's perspective

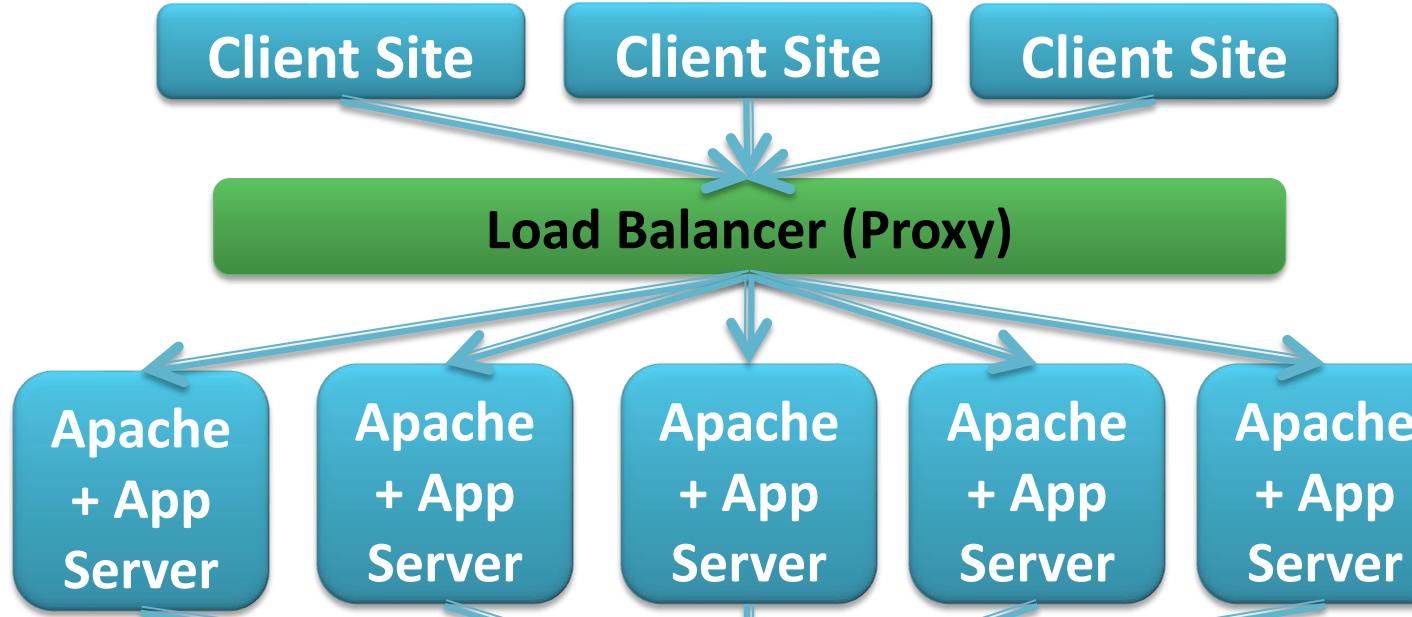
Scaling in the Cloud



Scaling in the Cloud



Scaling in the Cloud



Scalable and Elastic,
but limited consistency and
operational flexibility



ndra

BLOG Wisdom

- “If you want vast, on-demand scalability, you need a non-relational database.” Since scalability requirements:
 - Can change very quickly and,
 - Can grow very rapidly.
- Difficult to manage with a single in-house RDBMS server.
- RDBMS scale well:
 - When limited to a single node, **but**
 - Overwhelming complexity to scale on multiple server nodes.

The “NoSQL” movement

- Initially used for: “*Open-Source relational database that did not expose SQL interface*”
- Popularly used for: “*non-relational, distributed data stores that often did not attempt to provide ACID guarantees*”
- Gained widespread popularity through a number of open source projects
 - HBase, Cassandra, Voldemort, MongoDB, ...
- Scale-out, elasticity, flexible data model, high availability

NoSQL has no relation with SQL

- Micheal Stonebraker [[CACM Blog](#)]

- Term heavily used (and abused)
- Scalability and performance bottleneck not inherent to SQL
 - Scale-out, auto-partitioning, self-manageability can be achieved with SQL
- Different implementations of SQL engine for different application needs
- SQL provides flexibility, portability

No-SQL → Not Only SQL

- Recently renamed
- Encompass a broad category of “structured” storage solutions
 - RDBMS is a subset
 - Key Value stores
 - Document stores
 - Graph database
- The debate on appropriate characterization continues

Cloud Data Management Desiderata

SQL or NoSQL

- Scalability
- Elasticity
- Fault tolerance
- Self Manageability
- **Sacrifice consistency?**

Perspectives (James Hamilton)

I love **eventual consistency** but there are some applications that are much easier to implement with strong consistency. Many like eventual consistency because it allows us to scale-out nearly without bound *but it does come with a cost in programming model complexity.*



February 24, 2010

**Quest for Scalable, Fault-tolerant,
and Consistent Data Management in
the Cloud that provides
ELASTICITY**

Outline

- Data in the Cloud
- Data Platforms for Large Applications
 - Key value Stores
 - Transactional support in the cloud
- Multitenant Data Platforms
- Concluding Remarks

Key Value Stores

- Key-Valued data model
 - Key is the unique identifier
 - Key is the granularity for consistent access
 - Value can be structured or unstructured
- Gained widespread popularity
 - In house: **Bigtable** (Google), **PNUTS** (Yahoo!), **Dynamo** (Amazon)
 - Open source: **HBase**, **Hypertable**, **Cassandra**, **Voldemort**
- Popular choice for the modern breed of web-applications

Important Design Goals

- **Scale out: designed for scale**
 - Commodity hardware
 - Low latency updates
 - Sustain high update/insert throughput
- **Elasticity – scale up and down with load**
- **High availability – downtime implies lost revenue**
 - Replication (with multi-mastering)
 - Geographic replication
 - Automated failure recovery

Lower Priorities

- **No Complex querying functionality**
 - No support for SQL
 - CRUD operations through database specific API
 - **No support for joins**
 - Materialize simple join results in the relevant row
 - Give up normalization of data?
 - **No support for transactions**
 - Most data stores support single row transactions
 - Tunable consistency and availability (e.g., Dynamo)
- **Achieve high scalability**

Interplay with CAP

- Consistency, Availability, and Network Partitions
 - Only have two of the three together
- Large scale operations – be prepared for network partitions
- Role of CAP – During a network partition, choose between Consistency and Availability
 - RDBMS choose *consistency*
 - Key Value stores choose *availability* [low replica consistency]

Why sacrifice Consistency?

- It is a simple solution
 - nobody understands what sacrificing P means
 - sacrificing A is unacceptable in the Web
 - possible to push the problem to app developer
- C not needed in many applications
 - Banks do not implement ACID (classic example wrong)
 - Airline reservation only transacts reads (Huh?)
 - MySQL et al. ship by default in lower isolation level
- Data is noisy and inconsistent anyway
 - making it, say, 1% worse does not matter

C and A: In a Network Partition

- Dynamo – quorum based replication
 - Multi-mastering keys – Eventual Consistency
 - Tunable read and write quorums
 - Larger quorums – higher consistency, lower availability
 - Vector clocks to allow application supported reconciliation
- PNUTS – log based replication
 - Similar to log replay – reliable log multicast
 - Per record mastering – timeline consistency
 - Major outage might result in losing the tail of the log

**Too many choices – Which
system should I use?**

Benchmarking Serving Systems

[Cooper et al., SOCC 2010]

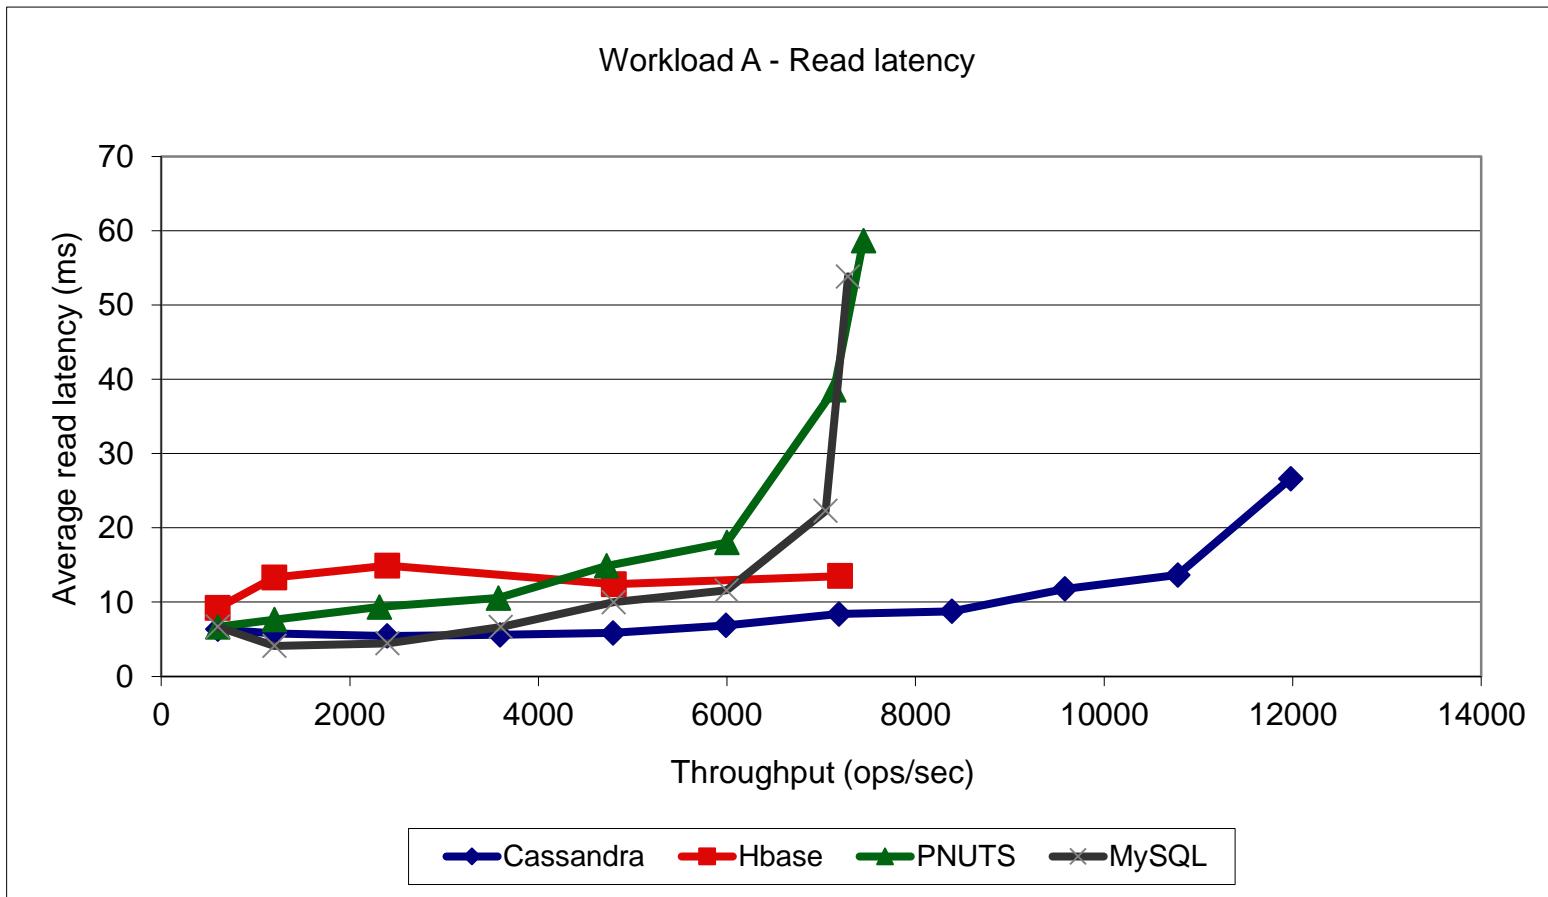
- A standard benchmarking tool for evaluating Key Value stores
- Evaluate different systems on common workloads
- Focus on performance and scale out

Benchmark tiers

- Tier 1 – Performance
 - Latency versus throughput as throughput increases
- Tier 2 – Scalability
 - Latency as database, system size increases
 - “Scale-out”
 - Latency as we elastically add servers
 - “Elastic speedup”

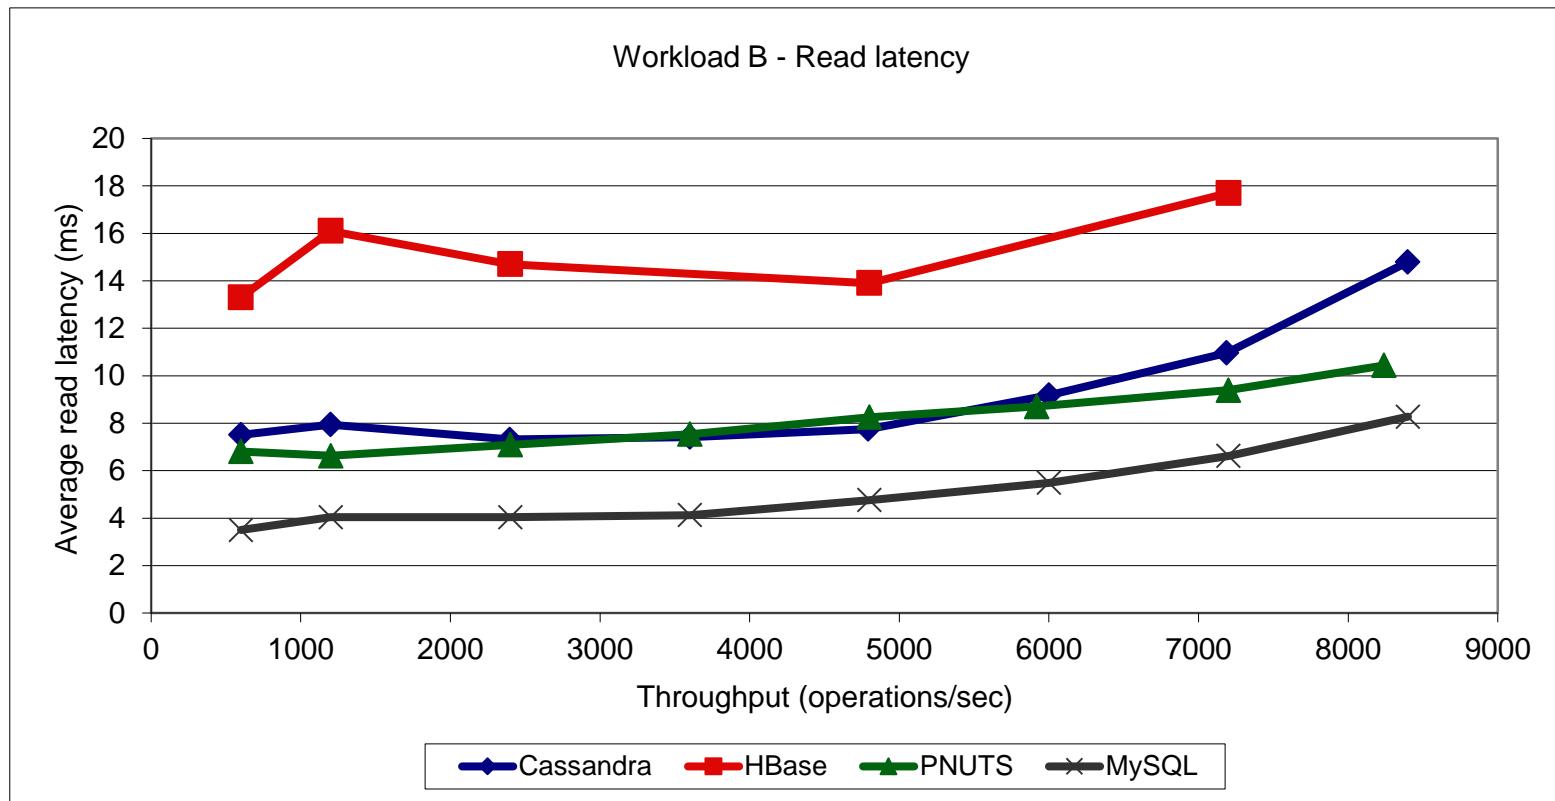
Workload A – Update heavy

- 50/50 Read/update



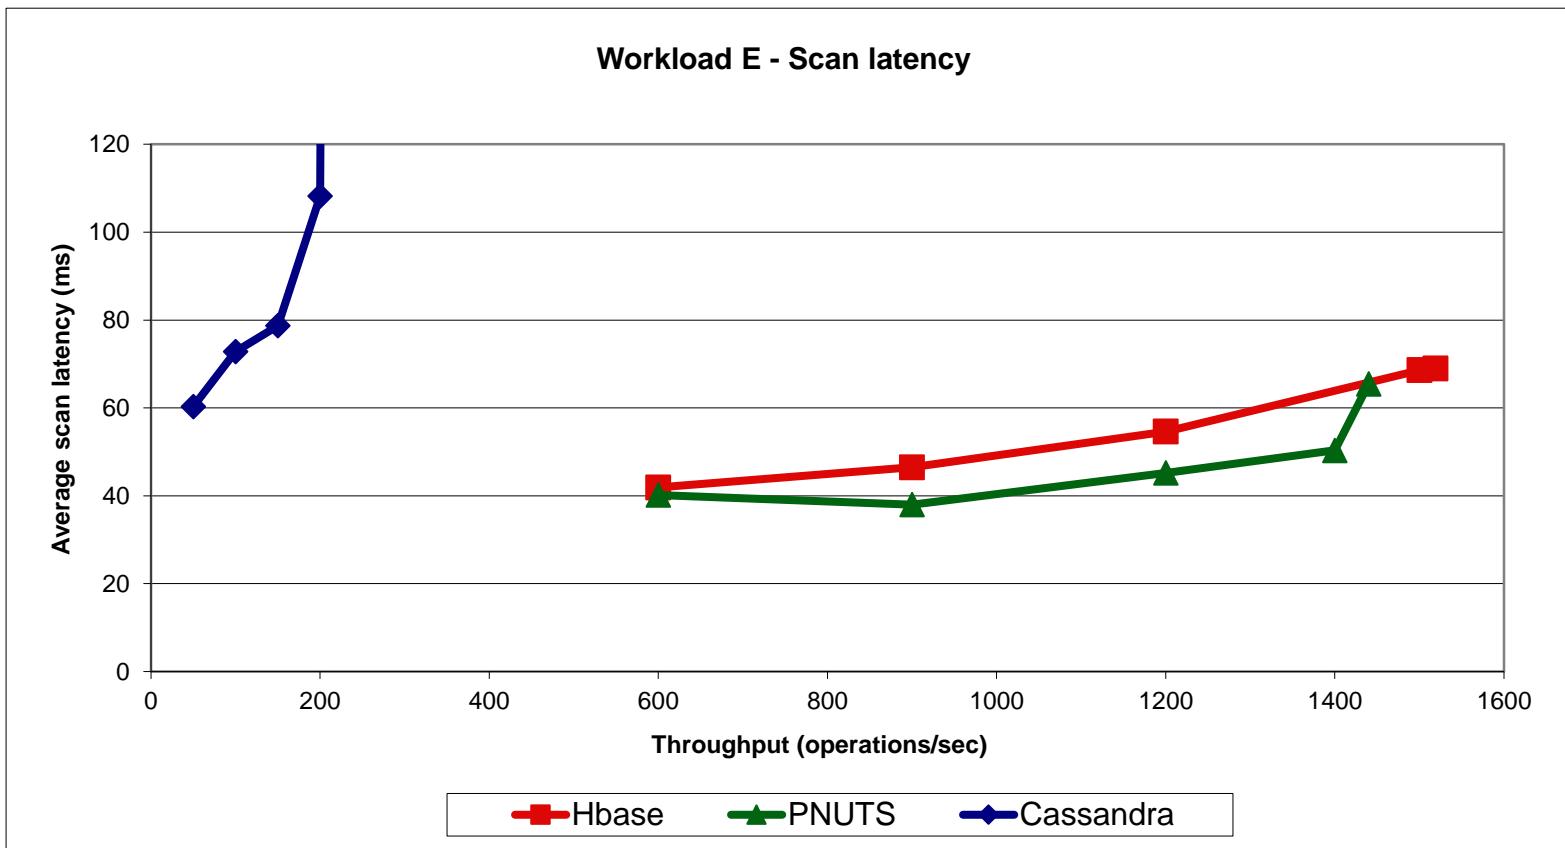
Workload B – Read heavy

- 95/5 Read/update



Workload E – short scans

- Scans of 1-100 records of size 1KB



Summary

- Different databases suitable for different workloads
- Evolving systems – landscape changing dramatically
- Active development community around open source systems